

Phi_K basic tutorial

This notebook guides you through the basic functionality of the phik package. The package offers functionality on three related topics:

1. Phik correlation matrix
2. Significance matrix
3. Outlier significance matrix

For more information on the underlying theory, the user is referred to our paper.

```
In [1]: %%capture
# install phik (if not installed yet)
import sys

!"{sys.executable}" -m pip install phik
```

```
In [2]: # import standard packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import itertools

import phik

from phik import resources
from phik.binning import bin_data
from phik.report import plot_correlation_matrix

%matplotlib inline
```

```
In [3]: # if one changes something in the phik-package one can automatically reload the package or module
%load_ext autoreload
%autoreload 2
```

Load data

A simulated dataset is part of the phik-package. The dataset concerns fake car insurance data. Load the dataset here:

```
In [4]: data = pd.read_csv( resources.fixture('fake_insurance_data.csv.gz') )
```

Take a first look at the data

Let's use a simple `data.head()` to get an idea of what the data looks like and inspect the different types of variables.

```
In [5]: data.head()
```

Out [5]:

	car_color	driver_age	area	mileage	car_size
0	black	26.377219	suburbs	156806.288398	XXL
1	black	58.976840	suburbs	74400.323559	XL
2	multicolor	55.744988	downtown	267856.748015	XXL
3	metallic	57.629139	downtown	259028.249060	XXL
4	green	21.490637	downtown	110712.216080	XL

Specify bin types

The phik-package offers a way to calculate correlations between variables of mixed types. Variable types can be inferred automatically although we recommend variable types to be specified by the user.

Because interval type variables need to be binned in order to calculate phik and the significance, a list of interval variables is created.

```
In [6]: data_types = {'severity': 'interval',
                    'driver_age': 'interval',
                    'satisfaction': 'ordinal',
                    'mileage': 'interval',
                    'car_size': 'ordinal',
                    'car_use': 'ordinal',
                    'car_color': 'categorical',
                    'area': 'categorical'}

interval_cols = [col for col, v in data_types.items() if v=='interval' and col in data.columns]
interval_cols
```

Out [6]: ['driver_age', 'mileage']

Visually inspect pairwise correlations

Bin the interval variables

To get a feeling for the data, let's bin the interval variables and create 2d histograms to inspect the correlations between variables. By binning the interval variables we can treat all variable types in the same way.

```
In [7]: # bin the interval variables
data_binned, binning_dict = bin_data(data, cols=interval_cols, retbins=True)
```

```
In [8]: binning_dict
```

```
Out[8]: {'driver_age': [(np.float64(18.030700817879502),
    np.float64(28.345238641214884)),
    (np.float64(28.345238641214884), np.float64(38.65977646455027)),
    (np.float64(38.65977646455027), np.float64(48.97431428788565)),
    (np.float64(48.97431428788565), np.float64(59.28885211122103)),
    (np.float64(59.28885211122103), np.float64(69.60338993455642)),
    (np.float64(69.60338993455642), np.float64(79.9179277578918)),
    (np.float64(79.9179277578918), np.float64(90.23246558122719)),
    (np.float64(90.23246558122719), np.float64(100.54700340456257)),
    (np.float64(100.54700340456257), np.float64(110.86154122789796)),
    (np.float64(110.86154122789796), np.float64(121.17607905123334))],
'mileage': [(np.float64(53.543057084421605), np.float64(30047.010932306846)),
    (np.float64(30047.010932306846), np.float64(60040.47880752927)),
    (np.float64(60040.47880752927), np.float64(90033.94668275169)),
    (np.float64(90033.94668275169), np.float64(120027.41455797412)),
    (np.float64(120027.41455797412), np.float64(150020.88243319656)),
    (np.float64(150020.88243319656), np.float64(180014.35030841897)),
    (np.float64(180014.35030841897), np.float64(210007.81818364142)),
    (np.float64(210007.81818364142), np.float64(240001.28605886383)),
    (np.float64(240001.28605886383), np.float64(269994.75393408624)),
    (np.float64(269994.75393408624), np.float64(299988.2218093087)))]}
```

```
In [9]: # plot each variable pair
plt.rc('text', usetex=False)

n=0
for i in range(len(data.columns)):
    n=n+i

ncols=3
nrows=int(np.ceil(n/ncols))
fig, axes = plt.subplots(nrows, ncols, figsize=(15,4*nrows))
ndecimals = 0

for i, comb in enumerate(itertools.combinations(data_binned.columns.values, 2)):

    c = int(i%ncols)
    r = int((i-c)/ncols )

    # get data
    c0, c1 = comb
    datahist = data_binned.groupby([c0,c1])[c0].count().to_frame().unstack().fillna(0)
    datahist.columns = datahist.columns.droplevel()

    # plot data
    img = axes[r][c].pcolormesh(datahist.values, edgecolor='w', linewidth=1)

    # axis ticks and tick labels
    if c0 in binning_dict.keys():
        ylabels = ['{1:.{0}f}_{2:.{0}f}'.format(ndecimals, binning_dict[c0][i][0], binning_dict[c0][i][1])
            for i in range(len(binning_dict[c0]))]
    else:
        ylabels = datahist.index

    if c1 in binning_dict.keys():
        xlabels = ['{1:.{0}f}_{2:.{0}f}'.format(ndecimals, binning_dict[c1][i][0], binning_dict[c1][i][1])
            for i in range(len(binning_dict[c1]))]
```

```
else:
```

```
    xlabel = datahist.columns
```

```
    # axis labels
```

```
    axes[r][c].set_yticks(np.arange(len(ylabels)) + 0.5)
```

```
    axes[r][c].set_xticks(np.arange(len(xlabels)) + 0.5)
```

```
    axes[r][c].set_xticklabels(xlabels, rotation='vertical')
```










```
    axes[r][c].set_yticklabels(ylabels, rotation='horizontal')
```







```
    axes[r][c].set_xlabel(datahist.columns.name)
```

```
    axes[r][c].set_ylabel(datahist.index.name)
```

```
    axes[r][c].set_title('data')
```

```
plt.tight_layout()
```





Correlation: mileage vs car_size

From the above plots it seems like there might be an interesting correlation between mileage and car_size. Let's see what phik correlation is measured for this data.

```
In [10]: x, y = data[['mileage', 'car_size']].T.values

print('phik      = %.2f'%phik.phik_from_array(x, y, num_vars=['x']))
print('significance = %.2f'%phik.significance_from_array(x, y, num_vars=['x'])[1])

phik      = 0.77
significance = 49.28
```

Indeed there is a correlation between these variables and the correlation is also significant. To better understand the correlation, we can have a look at the significance of excesses and deficits in the 2-dimensional contingency table, so-called "outlier significances".

```
In [11]: phik.outlier_significance_from_array(x, y, num_vars=['x'])
```


Out[11]:

	y	L	M	S	XL	XS	XXL
	53.5_30047.0	6.882155	21.483476	18.076204	-10.367152	10.820863	-22.423985
	30047.0_60040.5	20.034528	-0.251737	-3.408409	2.534277	-1.973628	-16.087822
	60040.5_90033.9	1.627610	-3.043497	-2.265809	10.215936	-1.246784	-8.318068
	90033.9_120027.4	-3.711579	-3.827278	-2.885475	12.999048	-1.638288	-7.185622
	120027.4_150020.9	-7.665861	-6.173001	-4.746762	9.629145	-2.841508	-0.504521
	150020.9_180014.4	-7.533189	-6.063786	-4.660049	1.559370	-2.785049	6.765549
	180014.4_210007.8	-5.541940	-4.425929	-3.360023	-4.802787	-1.942469	10.520540
	210007.8_240001.3	-3.496905	-2.745103	-2.030802	-5.850529	-1.100873	8.723925
	240001.3_269994.8	-5.275976	-4.207164	-3.186534	-8.616464	-1.830944	13.303101
	269994.8_299988.2	-8.014016	-6.458253	-4.973240	-12.868389	-2.989055	20.992824

The values displayed in the matrix are the significances of the outlier frequencies, i.e. a large value means that the measured frequency for that bin is significantly different from the expected frequency in that bin.

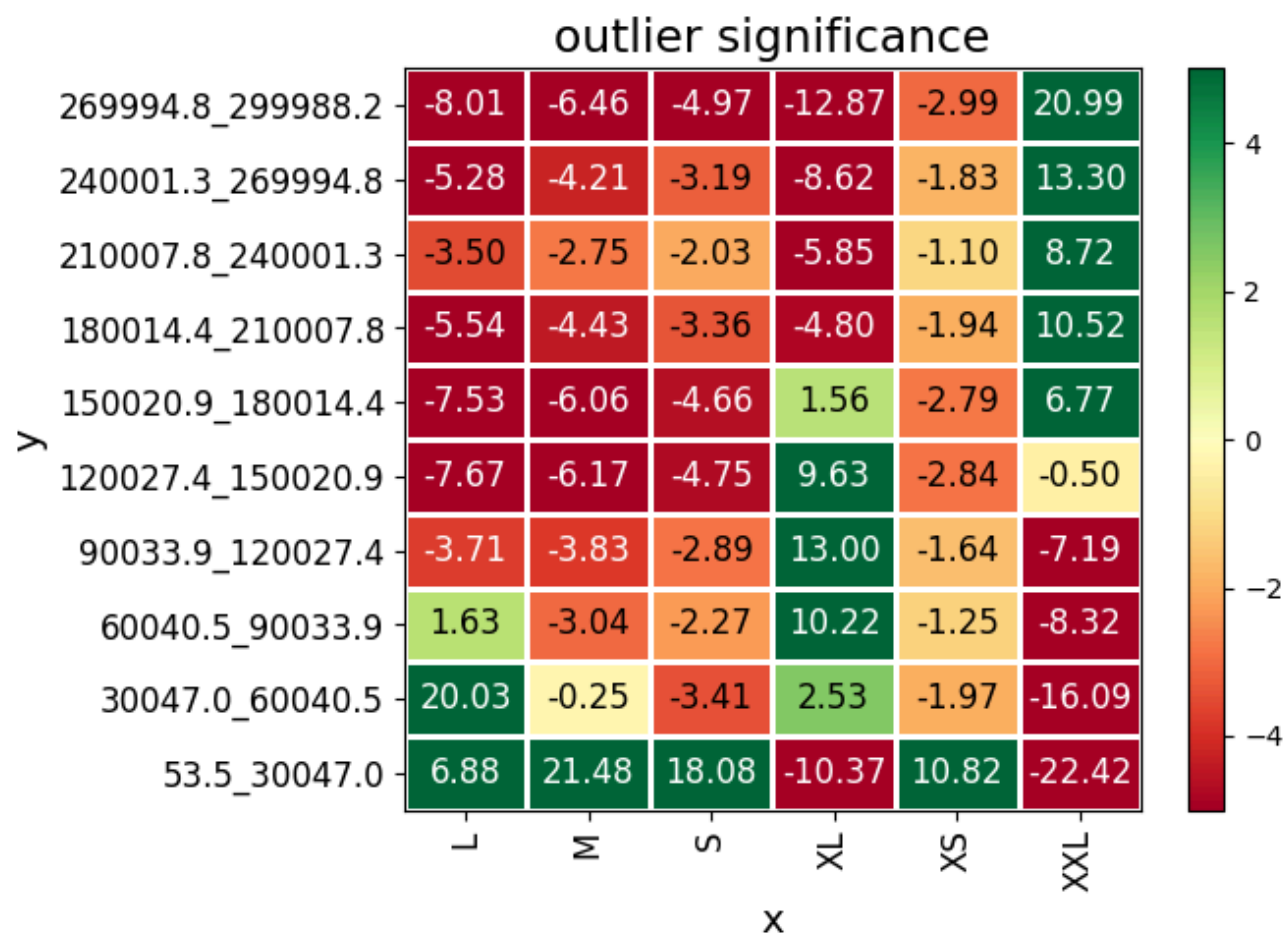
Let's visualise for easier interpretation.

In [12]:

```
outlier_signifs = phik.outlier_significance_from_array(x, y, num_vars=['x'])

zvalues = outlier_signifs.values
xlabels = outlier_signifs.columns
ylabels = outlier_signifs.index
xlabel = 'x'
ylabel = 'y'

plot_correlation_matrix(zvalues, x_labels=xlabels, y_labels=ylabels,
                        x_label=xlabel,y_label=ylabel,
                        vmin=-5, vmax=5, title='outlier significance',
                        identity_layout=False, fontsize_factor=1.2)
```



ϕ_k functions for dataframes

In our data we have 5 different columns, meaning we have to evaluate $4+3+2+1=10$ pairs of variables for possible correlations. In a large dataset, with many different variables, this can easily become a cumbersome task. Can we do this more efficient? yes! We have provided functions that work on dataframes, to allow you to calculate the phik correlation, significance and outlier significance for all different variable combinations at once.

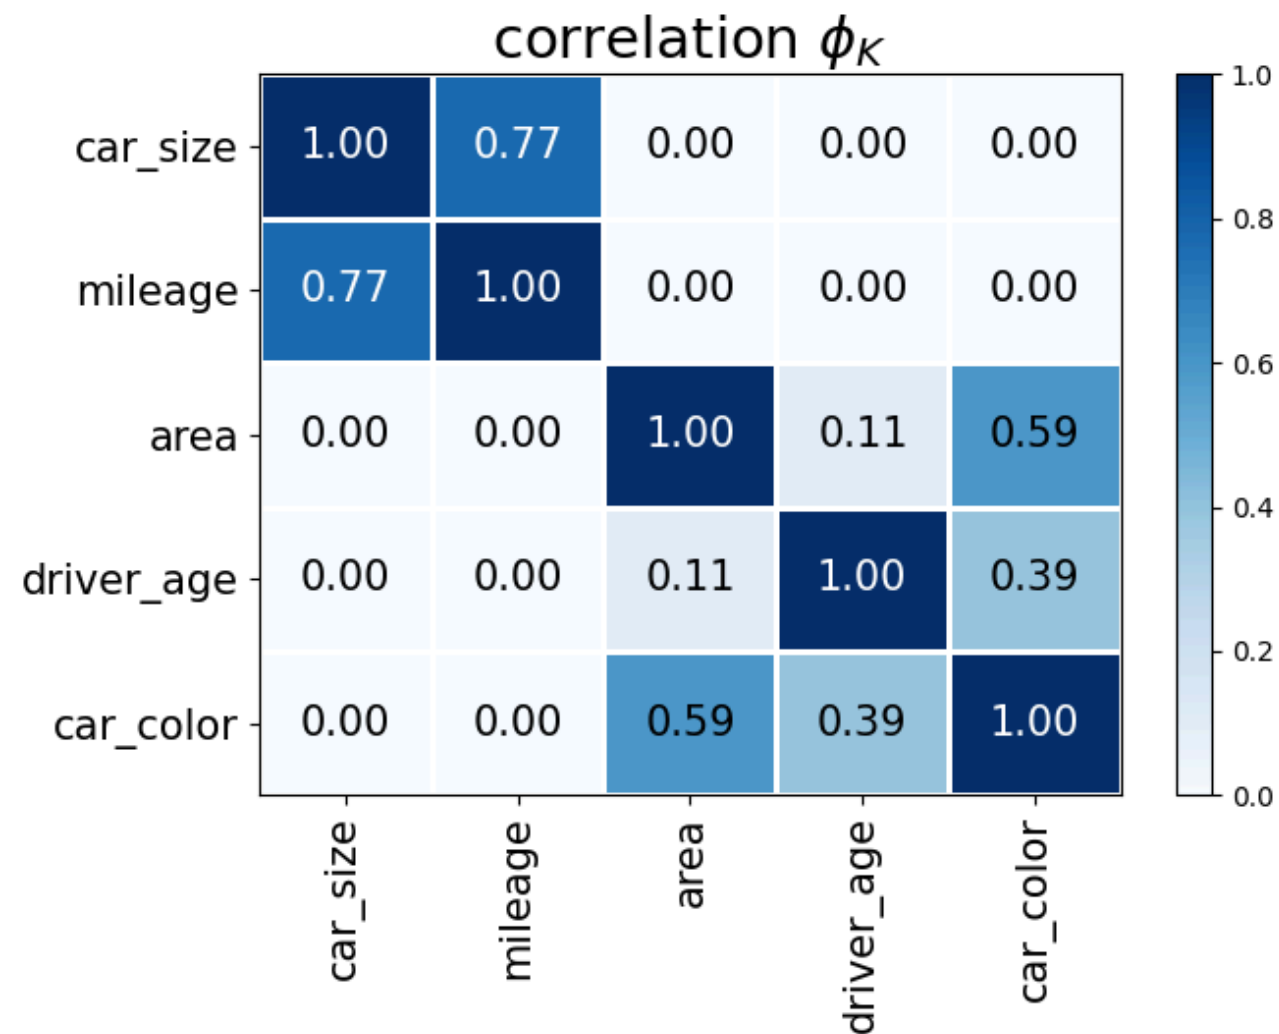
The functions are by default available after import of the phik package.

ϕ_k correlation matrix

Now let's start calculating the phik correlation coefficient between pairs of variables.

Note that the original dataset is used as input, the binning of interval variables is done automatically.

```
In [13]: phik_overview = data.phik_matrix(interval_cols=interval_cols)
phik_overview
```

Global correlation: ϕ_K

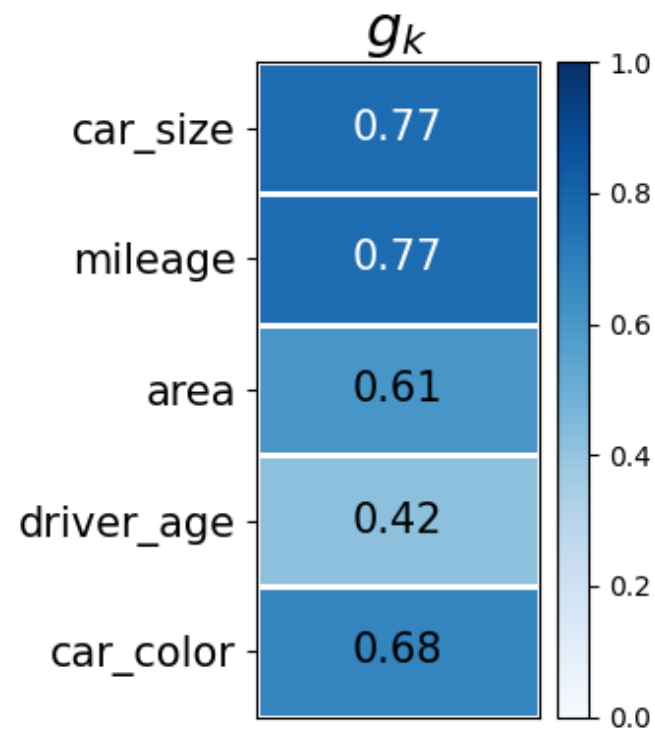
The global correlation coefficient is a measure of the total correlation of one variable to all other variables in the dataset. They give an indication of how well one variable can be modelled in terms of the other variables. A calculation of the global correlation coefficient is provided within the phik package.

```
In [16]: global_correlation, global_labels = data.global_phik(interval_cols=interval_cols)
for c, l in zip(global_correlation, global_labels):
    print(l, c[0])
```

```
car_color 0.6760316805291298
driver_age 0.41913014836282625
area 0.6057526929029953
mileage 0.7685888294891438
car_size 0.7685888294891438
```

```
In [17]: plot_correlation_matrix(global_correlation, x_labels=[''], y_labels=global_labels,
                                vmin=0, vmax=1, figsize=(3.5,4),
                                color_map='Blues', title=r'$\phi_K$',
                                fontsize_factor=1.5)

plt.tight_layout()
```



Statistical significance of the correlation: \$Z\$-score

When assessing correlations it is good practise to evaluate both the correlation and the significance of the correlation: a large correlation may be statistically insignificant, and vice versa a small correlation may be very significant. For instance, `scipy.stats.pearsonr` returns both the pearson correlation and the p-value. Similarly, the `phik` package offers functionality the calculate a significance matrix. Significance is defined as:

$$Z = \Phi^{-1}(1-p) \quad \Phi(z) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^z e^{-t^2/2} dt$$

Several corrections to the 'standard' p-value calculation are taken into account, making the method more robust for low statistics and sparse data cases. The user is referred to our paper for more details.

As a result, the calculation may take a few seconds.

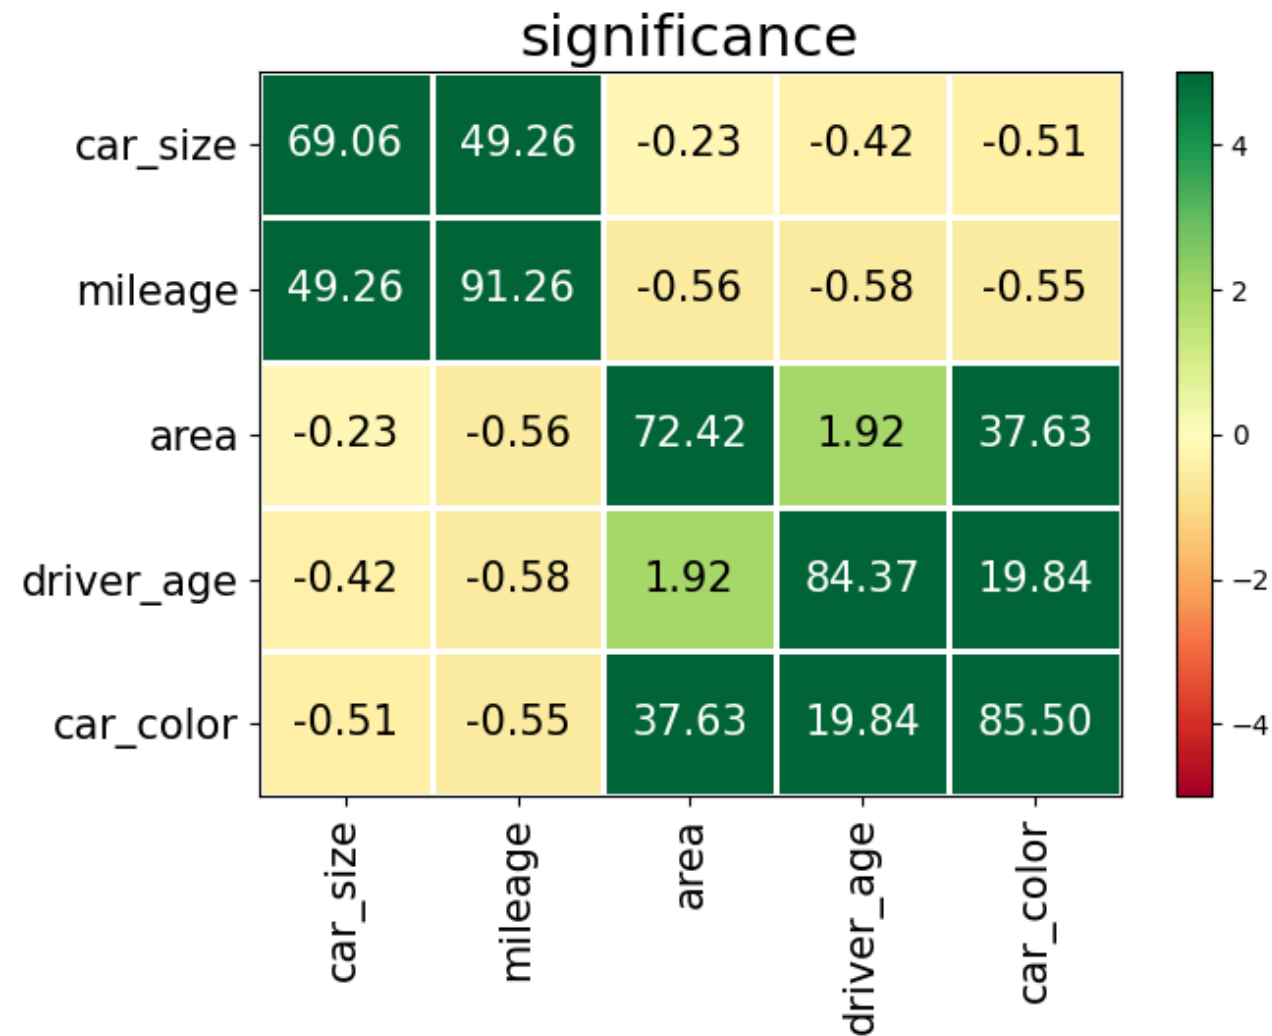
```
In [18]: significance_overview = data.significance_matrix(interval_cols=interval_cols)
significance_overview
```

Out[18]:

	car_color	driver_age	area	mileage	car_size
car_color	85.497799	19.844068	37.632867	-0.550590	-0.506238
driver_age	19.844068	84.369312	1.915829	-0.581718	-0.421274
area	37.632867	1.915829	72.416664	-0.562169	-0.226059
mileage	-0.550590	-0.581718	-0.562169	91.256826	49.260752
car_size	-0.506238	-0.421274	-0.226059	49.260752	69.059706

```
In [19]: plot_correlation_matrix(significance_overview.fillna(0).values, x_labels=significance_overview.columns,
                                y_labels=significance_overview.index, vmin=-5, vmax=5, title='significance',
```

```
plt.tight_layout()
usetex=False, fontsize_factor=1.5, figsize=(7,5.5))
```



Outlier significance

The normal pearson correlation between two interval variables is easy to interpret. However, the phik correlation between two variables of mixed type is not always easy to interpret, especially when it concerns categorical variables. Therefore, functionality is provided to detect "outliers": excesses and deficits over the expected frequencies in the contingency table of two variables.

Example 1: car_color versus area

For the categorical variable pair car_color - area we measured:

$$\phi_k = 0.59, \quad \text{significance} = 37.6$$

Let's use the outlier significance functionality to gain a better understanding of the significance correlation between car color and area.

```
In [20]: c1 = 'car_color'
         c0 = 'area'
```

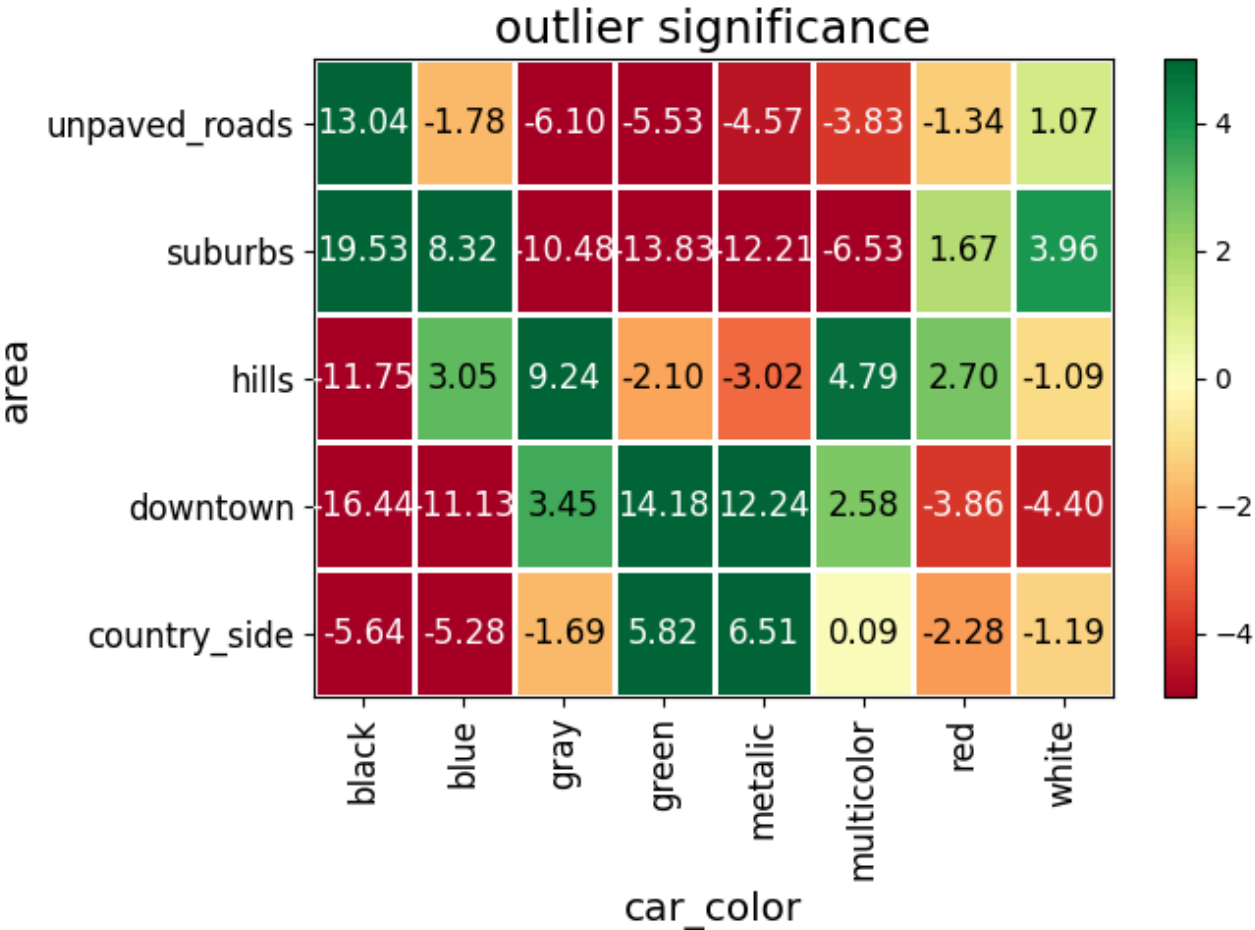
```
In [21]: outlier_signifs, binning_dict = data[[c0,c1]].outlier_significance_matrix(retbins=True)
outlier_signifs
```

Out[21]:

car_color	black	blue	gray	green	metallic	multicolor	red	white
area								
country_side	-5.636485	-5.277175	-1.688032	5.824664	6.506020	0.092630	-2.282290	-1.193468
downtown	-16.444319	-11.133951	3.450521	14.176441	12.238785	2.582260	-3.863796	-4.400847
hills	-11.752261	3.054306	9.244711	-2.097600	-3.022553	4.789665	2.695864	-1.092375
suburbs	19.533221	8.315558	-10.480108	-13.826572	-12.205322	-6.529846	1.667814	3.964555
unpaved_roads	13.043553	-1.776804	-6.095289	-5.532779	-4.566500	-3.829519	-1.339881	1.074386

```
In [22]: zvalues = outlier_signifs.values
xlabels = binning_dict[c1] if c1 in binning_dict.keys() else outlier_signifs.columns
ylabels = binning_dict[c0] if c0 in binning_dict.keys() else outlier_signifs.index
xlabel = c1
ylabel = c0

plot_correlation_matrix(zvalues, x_labels=xlabels, y_labels=ylabels,
                        x_label=xlabel,y_label=ylabel,
                        vmin=-5, vmax=5, title='outlier significance',
                        identity_layout=False, fontsize_factor=1.2)
```



The significance of each cell is expressed in terms of Z (one-sided).

Interesting, owners of a green car are more likely to live in the country side, and black cars are more likely to travel on unpaved roads!

Example 2: mileage versus car_size

For the categorical variable pair mileage - car_size we measured:

$\phi_k = 0.77$, $\mathrm{significance} = 46.3$

Let's use the outlier significance functionality to gain a better understanding of this significance correlation between mileage and car size.

```
In [23]: c0 = 'mileage'
c1 = 'car_size'

tmp_interval_cols = ['mileage']

In [24]: outlier_signifs, binning_dict = data[[c0,c1]].outlier_significance_matrix(interval_cols=tmp_interval_cols,
outlier_signifs
retbins=True)

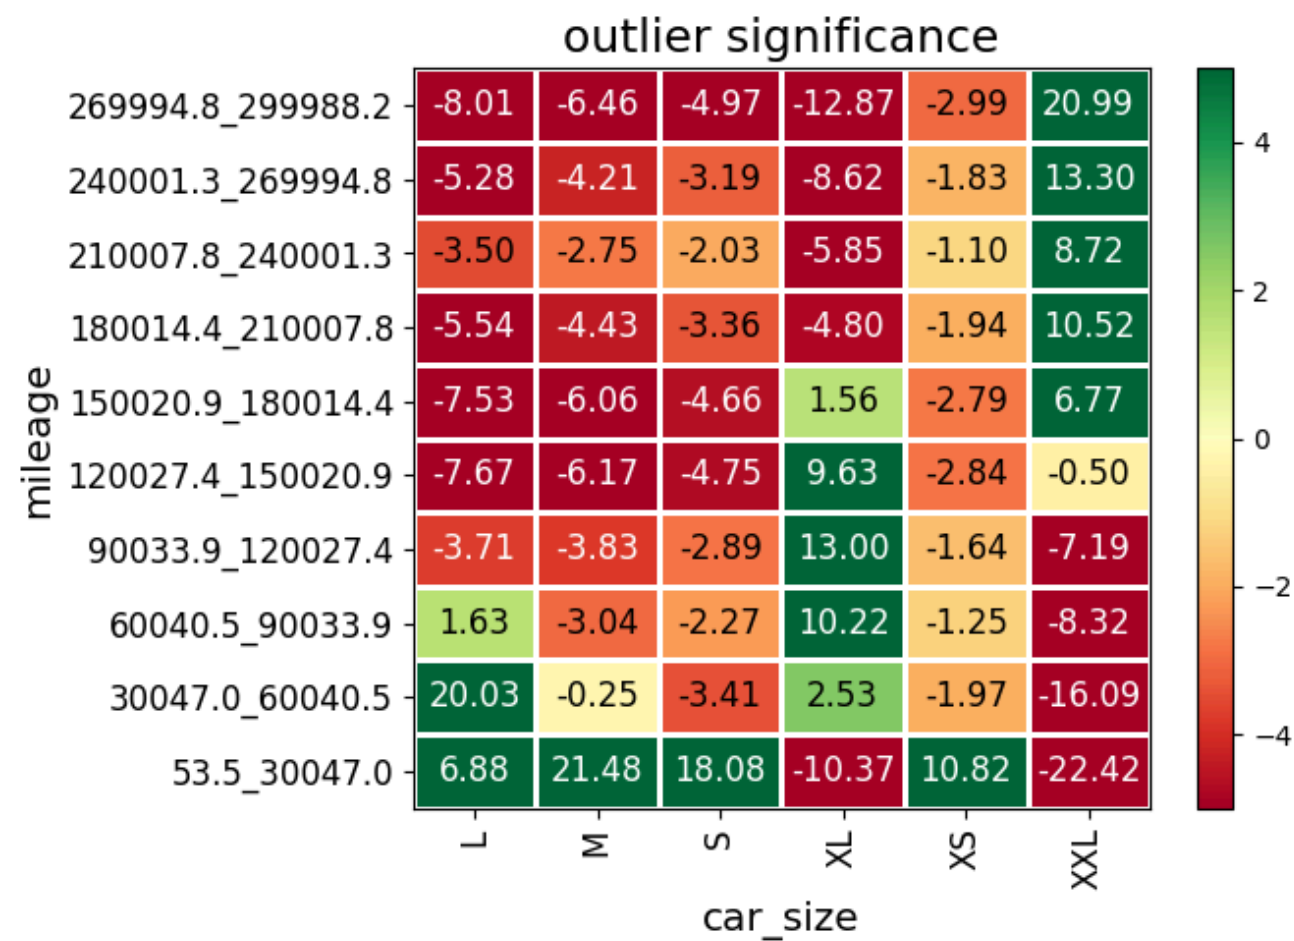
Out[24]:
```

	car_size	L	M	S	XL	XS	XXL
	53.5_30047.0	6.882155	21.483476	18.076204	-10.367152	10.820863	-22.423985
	30047.0_60040.5	20.034528	-0.251737	-3.408409	2.534277	-1.973628	-16.087822
	60040.5_90033.9	1.627610	-3.043497	-2.265809	10.215936	-1.246784	-8.318068
	90033.9_120027.4	-3.711579	-3.827278	-2.885475	12.999048	-1.638288	-7.185622
	120027.4_150020.9	-7.665861	-6.173001	-4.746762	9.629145	-2.841508	-0.504521
	150020.9_180014.4	-7.533189	-6.063786	-4.660049	1.559370	-2.785049	6.765549
	180014.4_210007.8	-5.541940	-4.425929	-3.360023	-4.802787	-1.942469	10.520540
	210007.8_240001.3	-3.496905	-2.745103	-2.030802	-5.850529	-1.100873	8.723925
	240001.3_269994.8	-5.275976	-4.207164	-3.186534	-8.616464	-1.830944	13.303101
	269994.8_299988.2	-8.014016	-6.458253	-4.973240	-12.868389	-2.989055	20.992824

Note that the interval variable mileage is binned automatically in 10 uniformly spaced bins!

```
In [25]: zvalues = outlier_signifs.values
xlabels = outlier_signifs.columns
ylabels = outlier_signifs.index
xlabel = c1
ylabel = c0

plot_correlation_matrix(zvalues, x_labels=xlabels, y_labels=ylabels,
x_label=xlabel,y_label=ylabel,
vmin=-5, vmax=5, title='outlier significance',
identity_layout=False, fontsize_factor=1.2)
```

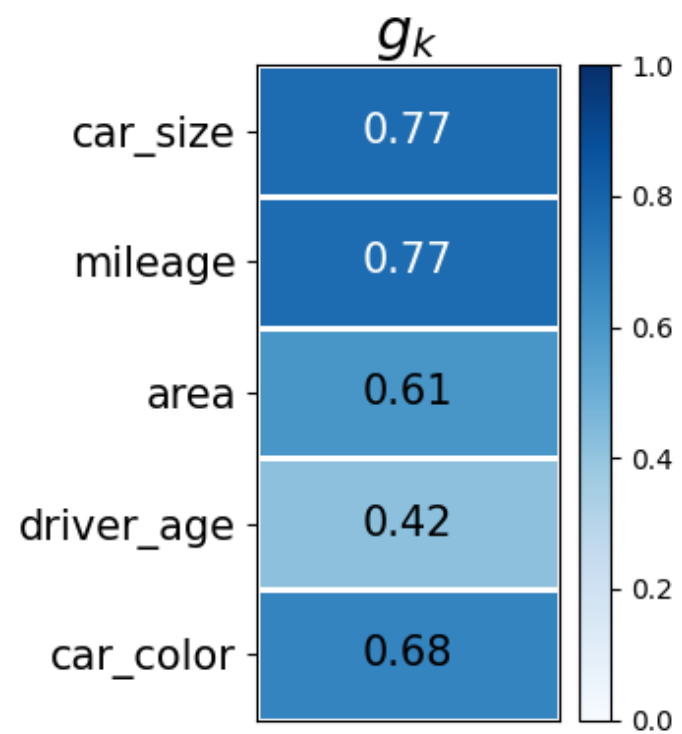
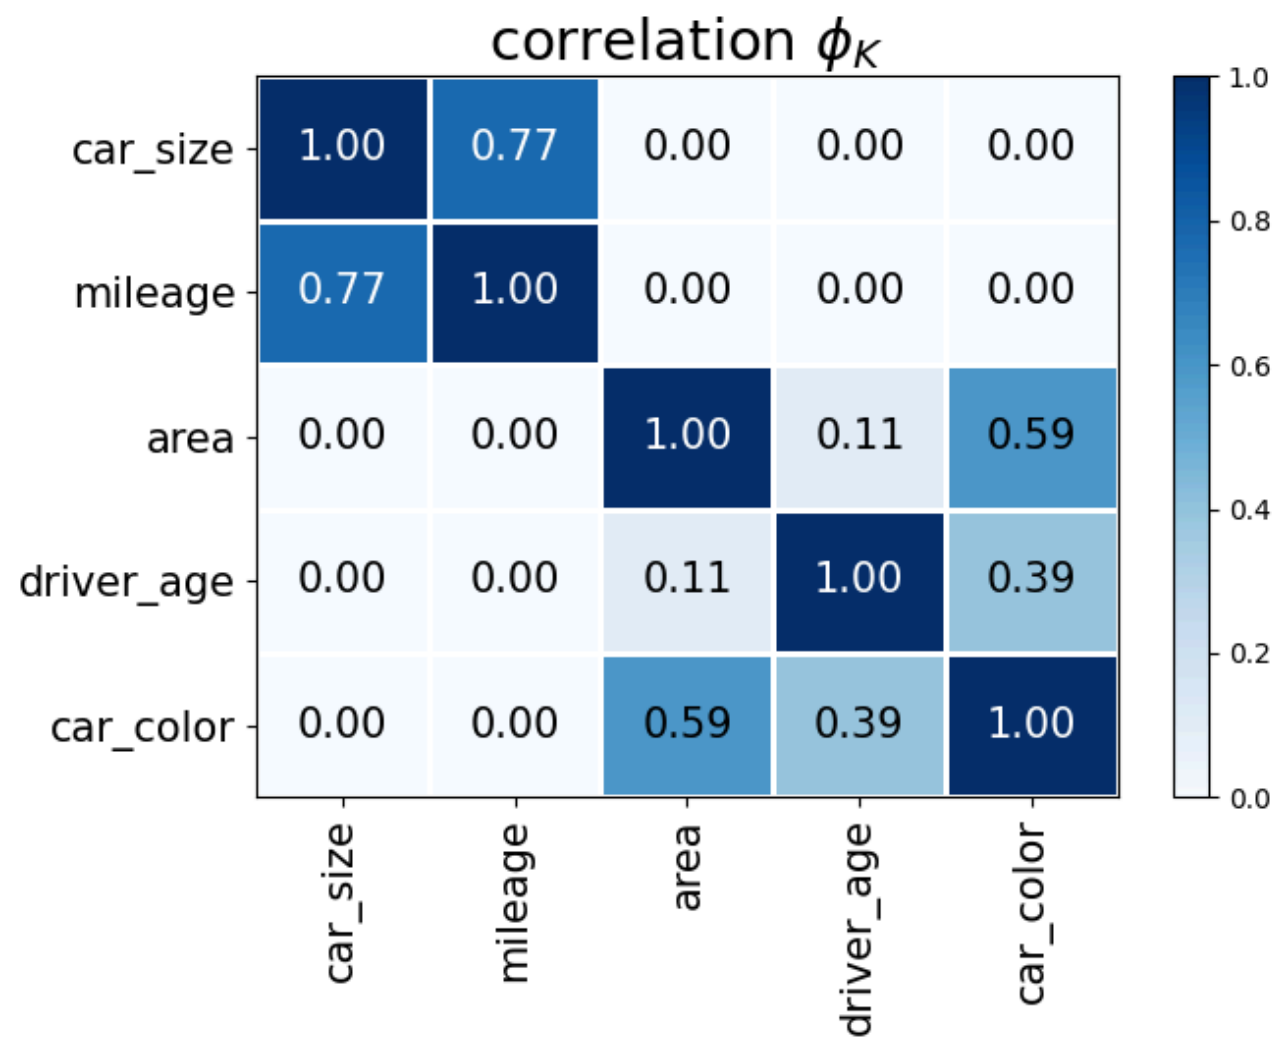
Correlation report

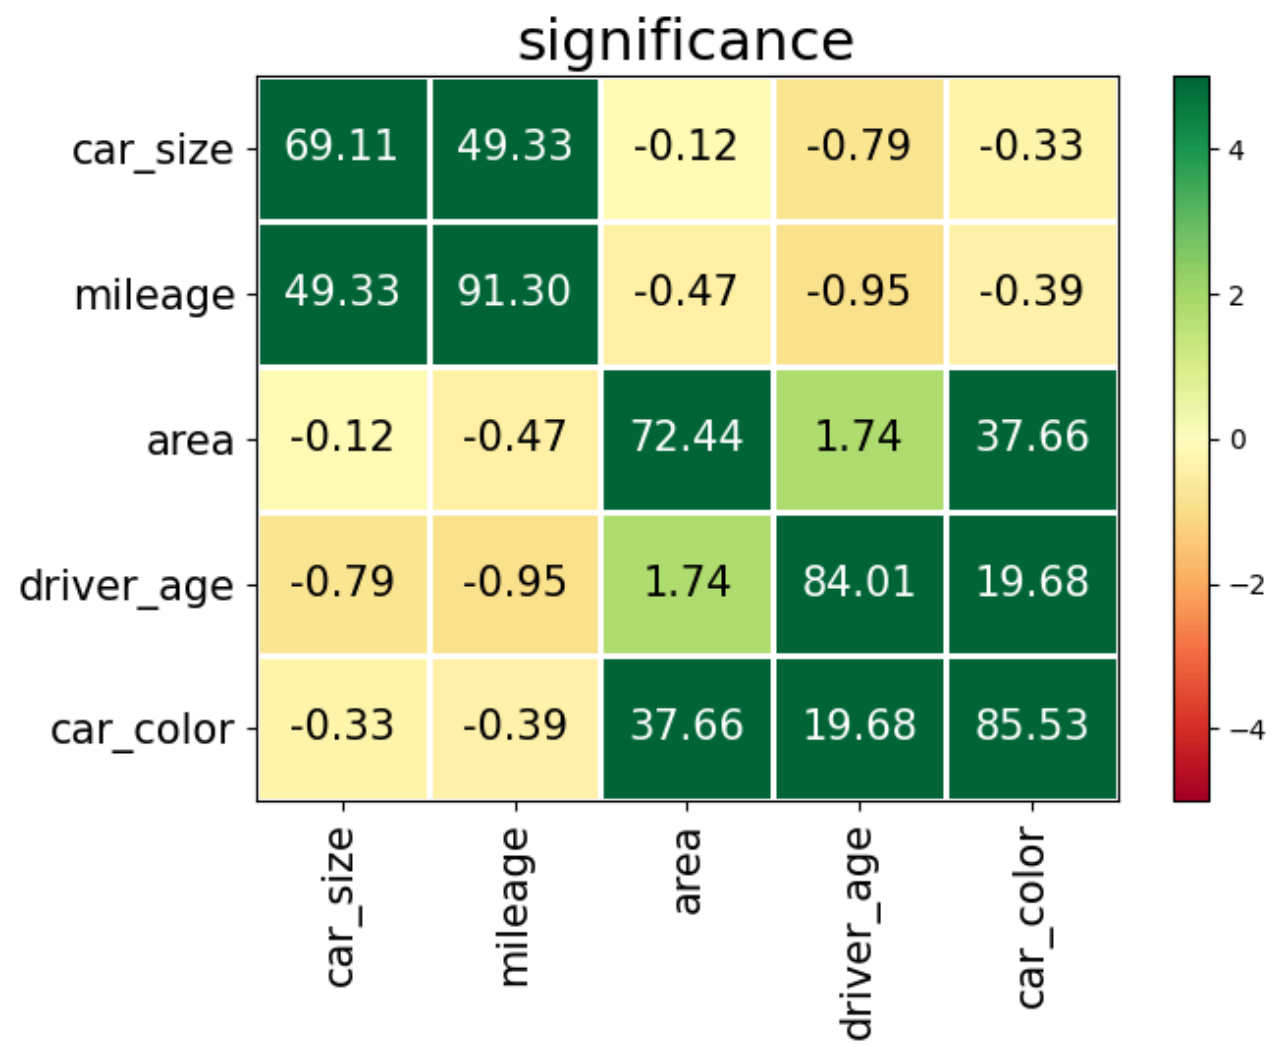
A full correlation report can be created automatically for a dataset by pairwise evaluation of all correlations, significances and outlier significances.

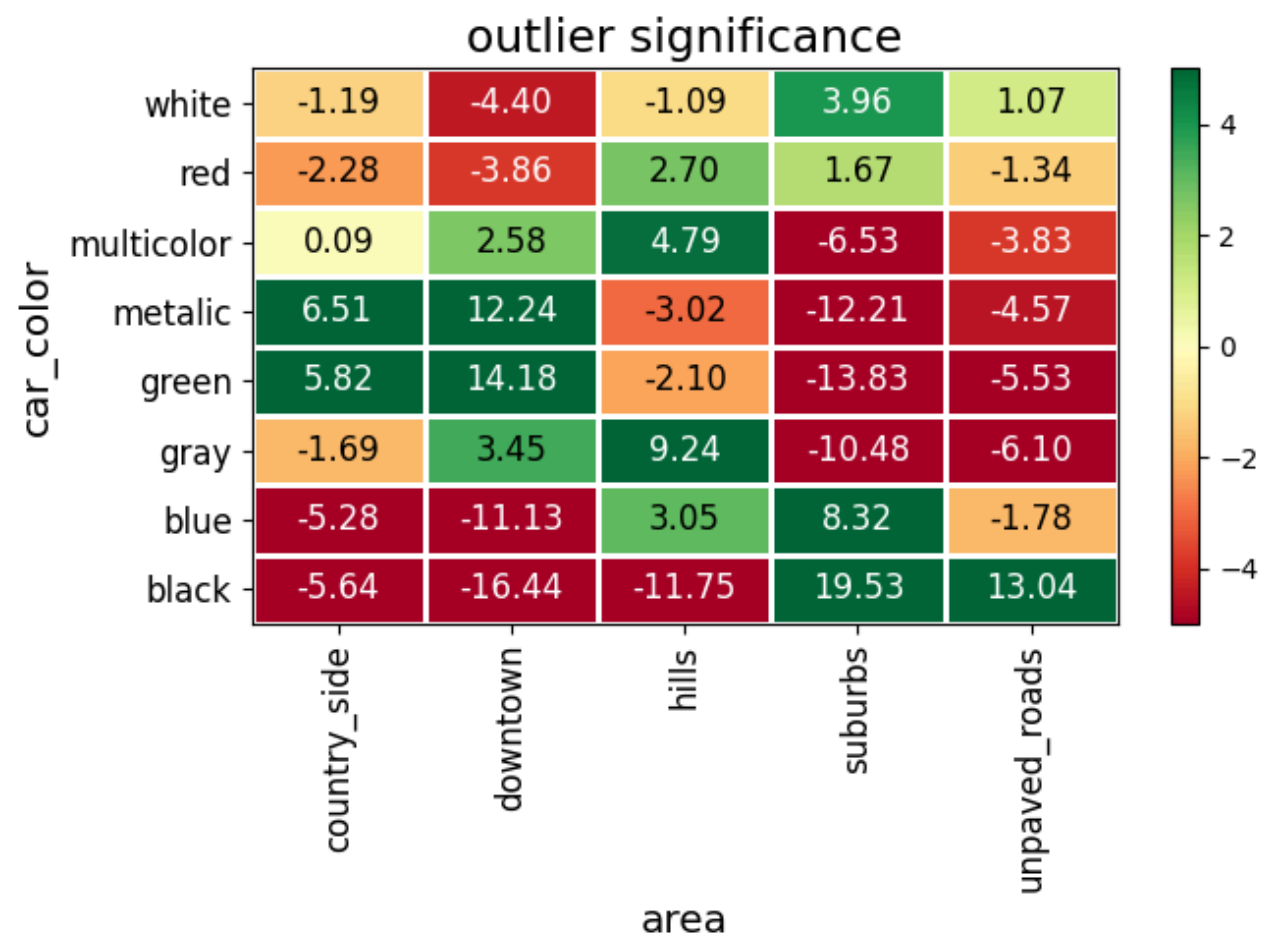
Note that for a dataset with many different columns the number of outlier significances plots can grow large very rapidly. Therefore, the feature is implemented to only evaluate outlier significances for those variable pairs with a significance and correlation larger than the given thresholds.

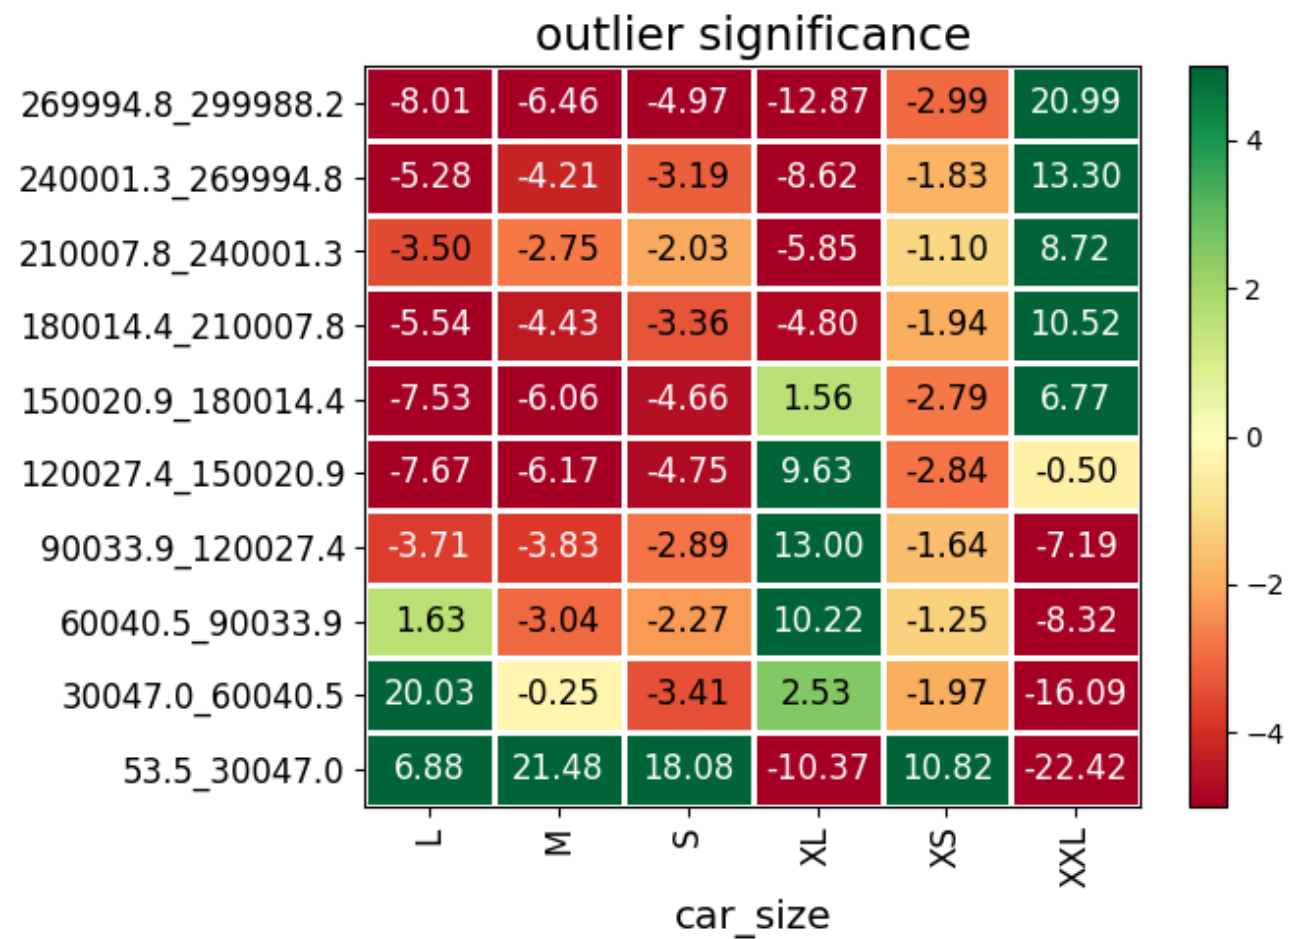
```
In [26]: from phik import report
         rep = report.correlation_report(data, significance_threshold=3, correlation_threshold=0.5)
```

interval columns not set, guessing: ['driver_age', 'mileage']









Recap

To summarize, the main functions in the phik correlation package working on a dataframe are:

- `df[twocols].hist2d()` or `series.hist2d(other_series)`
- `df.phik_matrix()`
- `df.global_phik()`
- `df.significance_matrix()`
- `df[twocols].outlier_significance_matrix()` or `series.hist2d(other_series)`
- `df.outlier_significance_matrices()`

```
In [27]: data[['driver_age', 'mileage']].hist2d()
# Alternatively: data['driver_age'].hist2d(data['mileage'])
```

interval columns not set, guessing: ['driver_age', 'mileage']

Out[27]:

	mileage	1	2	3	4	5	6	7	8	9	10
driver_age											
1	54.0	30.0	13.0	29.0	57.0	61.0	36.0	11.0	36.0	61.0	
2	51.0	41.0	13.0	17.0	64.0	41.0	35.0	12.0	30.0	75.0	
3	65.0	29.0	21.0	24.0	53.0	54.0	35.0	14.0	27.0	66.0	
4	49.0	23.0	17.0	30.0	50.0	54.0	21.0	17.0	25.0	53.0	
5	41.0	28.0	19.0	15.0	33.0	45.0	27.0	11.0	23.0	39.0	
6	21.0	12.0	5.0	12.0	26.0	25.0	10.0	8.0	10.0	15.0	
7	7.0	8.0	2.0	3.0	8.0	9.0	4.0	2.0	3.0	10.0	
8	5.0	2.0	0.0	1.0	6.0	1.0	1.0	1.0	1.0	3.0	
9	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	
10	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	

In [28]:

data.phik_matrix()

interval columns not set, guessing: ['driver_age', 'mileage']

Out[28]:

	car_color	driver_age	area	mileage	car_size
car_color	1.000000	0.389671	0.590456	0.000000	0.000000
driver_age	0.389671	1.000000	0.105506	0.000000	0.000000
area	0.590456	0.105506	1.000000	0.000000	0.000000
mileage	0.000000	0.000000	0.000000	1.000000	0.768589
car_size	0.000000	0.000000	0.000000	0.768589	1.000000

In [29]:

data.global_phik()

interval columns not set, guessing: ['driver_age', 'mileage']

Out[29]:

(array([[0.67603168], [0.41913015], [0.60575269], [0.76858883], [0.76858883]]), array(['car_color', 'driver_age', 'area', 'mileage', 'car_size'], dtype=object))
--

In [30]:

data.significance_matrix()

interval columns not set, guessing: ['driver_age', 'mileage']

```
Out[30]:
```

	car_color	driver_age	area	mileage	car_size
car_color	85.493401	19.868110	37.623334	-0.521680	-0.507110
driver_age	19.868110	84.377262	1.916490	-0.610847	-0.484669
area	37.623334	1.916490	72.418951	-0.546532	-0.187230
mileage	-0.521680	-0.610847	-0.546532	91.256984	49.261695
car_size	-0.507110	-0.484669	-0.187230	49.261695	69.063200

```
In [31]: data[['area', 'mileage']].outlier_significance_matrix()
interval columns not set, guessing: ['mileage']
Out[31]:
```

	53.5_30047.0	30047.0_60040.5	60040.5_90033.9	90033.9_120027.4	120027.4_150020.9	150020.9_180014.4	180014.4_210007.8	210007.8_240001.3	240001.3_269994.8	269994.8_290000.0
area										
country_side	-0.178204	0.192584	0.099073	-1.818924	0.103907	0.888225	0.251480	1.491532	0.016080	-0.000000
downtown	0.625586	-2.141106	0.541785	0.931452	-0.926177	0.582585	0.917301	-0.729204	-0.978248	-0.000000
hills	-0.826807	1.346029	-0.841090	0.008498	0.071356	1.028940	0.107051	0.080217	-0.397630	-0.000000
suburbs	0.972621	0.098196	0.712079	-0.059975	0.975536	-1.759184	-0.716170	-0.603623	0.059615	-0.000000
unpaved_roads	-1.469438	0.518041	-0.915745	-0.475005	-0.679357	-0.545607	-0.808321	0.901333	2.175620	-0.000000

```
In [32]: os_matrices = data.outlier_significance_matrices()
interval columns not set, guessing: ['driver_age', 'mileage']
In [33]: os_matrices.keys()
```

```
Out[33]: dict_keys(['car_color:driver_age', 'car_color:area', 'car_color:mileage', 'car_color:car_size', 'driver_age:area', 'driver_age:mileage', 'driver_age:car_size', 'area:mileage', 'area:car_size', 'mileage:car_size'])
In [34]: os_matrices['car_color:mileage']
```

Out[34]:

	53.5_30047.0	30047.0_60040.5	60040.5_90033.9	90033.9_120027.4	120027.4_150020.9	150020.9_180014.4	180014.4_210007.8	210007.8_240001.3	240001.3_269994.8	269994.8
--	--------------	-----------------	-----------------	------------------	-------------------	-------------------	-------------------	-------------------	-------------------	----------

car_color										
black	0.366425	1.387093	-1.320194	-1.274558	1.908010	-0.880578	-1.260293	-0.498607	-0.449794	
blue	-0.409760	-1.676495	-0.515491	1.969650	0.474818	-1.260466	-0.505987	-0.575699	1.675451	
gray	0.267888	-0.338430	-1.148378	0.048682	-0.238285	1.522769	1.282298	-0.812279	-0.589854	
green	-0.195906	0.399471	1.306739	-0.918723	0.373229	-0.859866	-0.862544	0.772096	-0.169912	
metallic	0.697184	-0.385806	0.672620	0.205599	-1.543500	1.168147	0.249471	0.513459	-0.174870	
multicolor	-0.478862	-0.591987	0.893783	0.050843	-2.104857	0.994621	0.109020	0.578461	-0.506460	
red	-0.551106	0.548949	1.370043	-0.050064	-1.822748	0.055521	2.208175	0.921955	0.219336	
white	-0.055025	1.219026	-0.596627	-0.427730	1.888206	-0.471175	-0.851352	-0.393054	-0.697834	