

```
In [1]: import os
import sys
sys.path.append(os.path.dirname(os.getcwd()))

from query_helper import get_connection, run_query

conn = get_connection()
```

Connection to database successful.

```
In [2]: """
All imports needed for EDA
"""

import pandas as pd
import matplotlib.pyplot as plt
```

```
In [3]: """
Load data into Python for analysis. Combine them.
"""

df = run_query(conn, """
SELECT DISTINCT
    patients.patient_id,
    patients.age,
    patients.gender,
    patients.city,
    patients.insurance_provider,
    patients.chronic_flag,
    patients.registration_date,

    visits.visit_id,
    visits.visit_date,
    visits.department,
    visits.visit_type,
    visits.length_of_stay_hours,
    visits.risk_score,
    visits.doctor_id,

    billing.bill_id,
    billing.billed_amount,
    billing.approved_amount,
    billing.claim_status,
    billing.payment_days,
    billing.billing_date

FROM
    patients
    INNER JOIN visits on (patients.patient_id = visits.patient_id)
    INNER JOIN billing on (visits.visit_id = billing.visit_id)
""", title="Load Combined Dataset"
)
```

=====

Load Combined Dataset

=====

📄 Query Plan:
SCAN billing USING INDEX idx_billing_date
SEARCH visits USING INTEGER PRIMARY KEY (rowid=?)
SEARCH patients USING INTEGER PRIMARY KEY (rowid=?)
📊 Results (25000 rows):

	patient_id	age	gender	city	insurance_provider	chronic_flag	registration_date	visit_id	visit_date	department	visit_type	length_of_stay_hours	risk_score	doctor_id	bill_id	billed_
0	2964	25	F	Chennai	HealthPlus	1	2025-07-04	3	2025-07-13	ICU	ER	34.36	Low	153	3	
1	4271	42	F	Mumbai	MediCareX	1	2025-10-31	669	2026-01-17	Neurology	ICU	12.74	Medium	138	669	2
2	1826	52	M	Pune	CareOne	1	2025-01-23	1264	2025-10-01	General	ICU	24.16	Low	123	1264	2
3	1224	52	M	Delhi	CareOne	1	2025-06-23	1326	2025-11-05	General	OPD	7.43	Low	134	1326	2
4	1318	39	M	Chennai	HealthPlus	1	2025-11-13	1455	2025-08-20	Cardiology	ER	17.72	High	136	1455	1
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
24995	4156	41	M	Bangalore	SecureLife	1	2025-06-23	22496	2025-10-11	Neurology	ICU	8.68	Medium	186	22496	2
24996	2166	26	M	Mumbai	HealthPlus	0	2025-05-14	24269	2025-05-20	General	ER	22.60	Low	180	24269	
24997	203	64	F	Pune	HealthPlus	1	2025-06-16	24461	2025-08-10	Orthopedics	ICU	36.48	Low	115	24461	
24998	1802	43	M	Delhi	CareOne	1	2025-04-14	24558	2025-09-27	ER	ICU	15.24	Low	141	24558	1
24999	2101	51	F	Mumbai	SecureLife	1	2025-10-14	24983	2025-04-25	Neurology	ICU	18.61	Low	167	24983	

25000 rows x 20 columns



```
In [4]: """
Analyze missing values in approved_amount, payment_days, and length_of_stay_hours.
"""

cols = ["approved_amount", "payment_days", "length_of_stay_hours"]

# Count and percentage of missing values
missing = df[cols].isnull().sum().rename("missing_count").to_frame()
missing["total"] = len(df)
missing["missing_pct"] = (missing["missing_count"] / missing["total"] * 100).round(2)
display(missing)

# Breakdown: are missing values concentrated in specific groups
```

```
for col in cols:
    if df[col].isnull().any():
        print(f"\n{'='*60}")
        print(f"  Missing '{col}' by department & claim_status")
        print(f"\n{'='*60}")
        breakdown = (
            df[df[col].isnull()]
            .groupby(["department", "claim_status"])
            .size()
            .reset_index(name="missing_count")
            .sort_values("missing_count", ascending=False)
        )
        display(breakdown)
```

	missing_count	total	missing_pct
approved_amount	1318	25000	5.27
payment_days	790	25000	3.16
length_of_stay_hours	0	25000	0.00

=====

Missing 'approved_amount' by department & claim_status

=====

	department	claim_status	missing_count
3	ER	Paid	140
15	Orthopedics	Paid	138
6	General	Paid	137
12	Neurology	Paid	137
0	Cardiology	Paid	136
9	ICU	Paid	129
4	ER	Pending	57
1	Cardiology	Pending	52
7	General	Pending	51
10	ICU	Pending	50
13	Neurology	Pending	46
16	Orthopedics	Pending	45
8	General	Rejected	36
2	Cardiology	Rejected	36
5	ER	Rejected	33
14	Neurology	Rejected	32
17	Orthopedics	Rejected	32
11	ICU	Rejected	31

=====			
Missing 'payment_days' by department & claim_status			
=====			
	department	claim_status	missing_count
0	Cardiology	Paid	95
6	General	Paid	86
12	Neurology	Paid	84
15	Orthopedics	Paid	67
9	ICU	Paid	67
3	ER	Paid	60
1	Cardiology	Pending	44
4	ER	Pending	40
16	Orthopedics	Pending	40
10	ICU	Pending	32
7	General	Pending	28
17	Orthopedics	Rejected	28
13	Neurology	Pending	24
14	Neurology	Rejected	23
11	ICU	Rejected	22
8	General	Rejected	18
2	Cardiology	Rejected	18
5	ER	Rejected	14

```
In [5]: """
Fix missing values for approved_amount column.
Fallback order: patient median → department median → global median.
"""

df["approved_amount"] = pd.to_numeric(df["approved_amount"], errors="coerce")

# Per-patient median
patient_median = df.groupby("patient_id")["approved_amount"].transform("median")

# Per-department median
dept_median = df.groupby("department")["approved_amount"].transform("median")

# Global median
global_median = df["approved_amount"].median()

df["approved_amount"] = (
    df["approved_amount"]
    .fillna(patient_median)
    .fillna(dept_median)
    .fillna(global_median)
)
```

```
        .fillna(dept_median)
        .fillna(global_median)
    )

print(f"Remaining missing approved_amount: {df['approved_amount'].isnull().sum()}")
```

Remaining missing approved_amount: 0

```
In [6]: """
Fix missing values for payment_days column.
Fallback order: patient median → department median → global median.
"""

df["payment_days"] = pd.to_numeric(df["payment_days"], errors="coerce")

# Per-patient median
patient_median = df.groupby("patient_id")["payment_days"].transform("median")

# Per-department median
dept_median = df.groupby("department")["payment_days"].transform("median")

# Global median
global_median = df["payment_days"].median()

df["payment_days"] = (
    df["payment_days"]
    .fillna(patient_median)
    .fillna(dept_median)
    .fillna(global_median)
)

print(f"Remaining missing payment_days: {df['payment_days'].isnull().sum()}")
```

Remaining missing payment_days: 0

```
In [7]: """
Fix missing values for length_of_stay_hours column.
Fallback order: patient median → department median → global median.
"""

df["length_of_stay_hours"] = pd.to_numeric(df["length_of_stay_hours"], errors="coerce")
# Per-patient median
patient_median = df.groupby("patient_id")["length_of_stay_hours"].transform("median")
# Per-department median
dept_median = df.groupby("department")["length_of_stay_hours"].transform("median")
# Global median
global_median = df["length_of_stay_hours"].median()
df["length_of_stay_hours"] = (
    df["length_of_stay_hours"]
    .fillna(patient_median)
    .fillna(dept_median)
    .fillna(global_median)
)
print(f"Remaining missing length_of_stay_hours: {df['length_of_stay_hours'].isnull().sum()}")
```

Remaining missing length_of_stay_hours: 0

```
In [8]: """
Print summary statistics for the cleaned columns.
"""
```

```
cols = ["approved_amount", "payment_days", "length_of_stay_hours"]

# Count and percentage of missing values
missing = df[cols].isnull().sum().rename("missing_count").to_frame()
missing["total"] = len(df)
missing["missing_pct"] = (missing["missing_count"] / missing["total"] * 100).round(2)
display(missing)
```

	missing_count	total	missing_pct
approved_amount	0	25000	0.0
payment_days	0	25000	0.0
length_of_stay_hours	0	25000	0.0

```
In [9]: """
Perform distribution analysis by department, visit_type, insurance_provider, and city.
"""

cat_cols = ["department", "visit_type", "insurance_provider", "city", "risk_score"]

fig, axes = plt.subplots(2, 2, figsize=(16, 10))
axes = axes.flatten()

for ax, col in zip(axes, cat_cols):
    counts = df[col].value_counts()
    counts.plot(kind="bar", ax=ax, edgecolor="black")
    ax.set_title(f"Distribution of {col}", fontsize=13)
    ax.set_ylabel("Count")
    ax.set_xlabel(col)
    ax.tick_params(axis="x", rotation=45)

    # Annotate bars
    for bar in ax.patches:
        ax.annotate(
            f"{int(bar.get_height())}",
            (bar.get_x() + bar.get_width() / 2, bar.get_height()),
            ha="center", va="bottom", fontsize=9,
        )

plt.tight_layout()
plt.show()





# Summary statistics for numeric columns split by each categorical column
numeric_cols = ["billed_amount", "approved_amount", "length_of_stay_hours", "payment_days"]

# Ensure numeric types (SQLite NULLs can cause string dtype)
for c in numeric_cols:
    df[c] = pd.to_numeric(df[c], errors="coerce")

for cat in cat_cols:
    print(f"\n{'='*60}")
    print(f" Numeric summary grouped by '{cat}'")
```



```
print(f"{'='*60}")
summary = df.groupby(cat)[numeric_cols].agg(["mean", "median", "std"]).round(2)
display(summary)
```



```
=====
Numeric summary grouped by 'department'
=====
```

	billed_amount			approved_amount			length_of_stay_hours			payment_days		
	mean	median	std	mean	median	std	mean	median	std	mean	median	std
department												
Cardiology	20695.18	19582.89	12404.78	16041.99	14047.26	13367.78	19.60	18.01	12.60	12.92	12.0	7.12
ER	21015.87	19686.06	12669.38	16383.82	14629.48	13588.52	19.53	18.12	12.27	13.13	13.0	7.27
General	20608.20	19600.60	12492.02	16109.87	14227.30	13392.48	19.43	18.13	12.30	12.96	13.0	7.01
ICU	20855.75	19650.86	12463.45	16369.09	14461.38	13472.09	19.36	17.85	12.30	12.94	12.0	7.12
Neurology	20962.80	19609.85	12760.59	16308.74	14516.12	13609.47	19.72	18.63	12.18	13.03	12.0	7.21
Orthopedics	21088.25	19754.84	12839.03	16461.46	14618.56	13821.96	19.66	18.38	12.23	13.24	13.0	7.25

```
=====
Numeric summary grouped by 'visit_type'
=====
```

		billed_amount			approved_amount			length_of_stay_hours			payment_days		
		mean	median	std	mean	median	std	mean	median	std	mean	median	std
visit_type													
	ER	20866.52	19699.01	12561.81	16127.37	14288.78	13569.30	19.41	18.08	12.40	12.99	12.0	7.19
	ICU	20770.36	19534.17	12583.20	16220.84	14362.41	13469.91	19.53	18.24	12.11	13.12	13.0	7.16
	OPD	20973.67	19757.80	12673.87	16486.76	14534.89	13585.69	19.71	18.28	12.42	13.00	13.0	7.14

```
=====
Numeric summary grouped by 'insurance_provider'
=====
```

		billed_amount			approved_amount			length_of_stay_hours			payment_days		
		mean	median	std	mean	median	std	mean	median	std	mean	median	std
insurance_provider													
	CareOne	20803.44	19579.55	12652.09	16208.49	14108.64	13556.96	19.54	18.42	12.23	13.01	12.0	7.16
	HealthPlus	20929.38	19647.82	12489.16	16338.95	14599.02	13373.94	19.85	18.60	12.43	13.06	13.0	7.16
	MediCareX	20604.89	19422.75	12606.57	16076.41	14354.05	13517.98	19.28	17.84	12.25	12.98	12.0	7.16
	SecureLife	21171.67	19941.92	12675.19	16511.13	14662.16	13726.78	19.55	17.97	12.34	13.09	13.0	7.17

```
=====
Numeric summary grouped by 'city'
=====
```


city	billed_amount			approved_amount			length_of_stay_hours			payment_days		
	mean	median	std	mean	median	std	mean	median	std	mean	median	std
Bangalore	21044.48	19718.10	12755.30	16328.54	14389.03	13607.74	19.57	18.26	12.20	12.97	12.0	7.22
Chennai	20778.88	19597.59	12560.50	16197.51	14364.75	13485.34	19.58	18.34	12.22	12.94	13.0	7.17
Delhi	20891.34	19558.40	12650.77	16398.56	14490.78	13596.35	19.47	18.23	12.27	13.15	12.0	7.35
Hyderabad	20846.17	19637.29	12570.46	16136.14	14176.64	13581.82	19.78	18.46	12.33	12.99	13.0	7.15
Mumbai	20962.58	19722.04	12605.56	16422.90	14642.40	13530.82	19.30	17.89	12.40	13.04	13.0	6.98
Pune	20699.97	19571.90	12498.36	16195.36	14355.29	13454.10	19.59	18.03	12.45	13.12	13.0	7.11

=====
Numeric summary grouped by 'risk_score'
=====

risk_score	billed_amount			approved_amount			length_of_stay_hours			payment_days		
	mean	median	std	mean	median	std	mean	median	std	mean	median	std
High	20702.15	19705.76	12484.61	16000.89	14140.96	13420.43	19.76	18.22	12.83	12.99	13.0	7.18
Low	20900.72	19631.70	12587.42	16361.32	14434.60	13534.22	19.15	18.08	11.50	13.06	13.0	7.12
Medium	20934.15	19634.69	12719.16	16327.65	14519.56	13636.81	20.08	18.52	13.22	13.03	13.0	7.22

In [10]:

```
"""
Detect and classify outliers in billed_amount, payment_days, and length_of_stay_hours.
"""

outlier_cols = ["billed_amount", "payment_days", "length_of_stay_hours"]
outlier_summary = []

for col in outlier_cols:
    plt.figure(figsize=(8, 4))
    plt.boxplot(df[col].dropna(), vert=False)
    plt.title(f"Boxplot of {col}")
    plt.xlabel(col)
    plt.show()


    # Calculate IQR
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Classify outliers (local variable only – not added to df)
    outlier_flag = df[col].apply(lambda x: "Outlier" if x < lower_bound or x > upper_bound else "Normal")
    print(f"\nOutlier classification for {col}:")
    display(pd.DataFrame({"count": outlier_flag.value_counts()}))

    # Collect counts for summary table
```

```
counts = outlier_flag.value_counts()
outlier_summary.append({
    "column": col,
    "normal": counts.get("Normal", 0),
    "outlier": counts.get("Outlier", 0),
    "outlier_pct": round(counts.get("Outlier", 0) / len(df) * 100, 2),
})


# Summary table: Normal vs Outlier per column
print(f"\n{'='*60}")
print("  Outlier Summary (Normal vs Outlier)")
print(f"{'='*60}")
outlier_summary_df = pd.DataFrame(outlier_summary).set_index("column")
display(outlier_summary_df)
```



Outlier classification for billed_amount:

count	
billed_amount	
Normal	24627
Outlier	373


Boxplot of payment_days



Outlier classification for payment_days:

count	
payment_days	
Normal	24489
Outlier	511

Boxplot of length_of_stay_hours



Outlier classification for length_of_stay_hours:

	count
length_of_stay_hours	
Normal	24744
Outlier	256

=====

Outlier Summary (Normal vs Outlier)

=====

	normal	outlier	outlier_pct
column			
billed_amount	24627	373	1.49
payment_days	24489	511	2.04
length_of_stay_hours	24744	256	1.02

```
In [11]: df.to_csv(os.path.join(os.path.dirname(os.getcwd()), "Database", "interim_eda.csv"), index=False)
print(f"Saved {len(df)} rows to Database/interim_eda.csv")
```

Saved 25000 rows to Database/interim_eda.csv