

Table of Contents

<i>Exercise notebook code and output</i>	2
<i>Documentation</i>	3
<i>Presentation</i>	20

Exercise notebook code and output

All work files, including report and notebook are located here in GitHub

<https://github.com/mramsundar/IITM-Graded->

Projects/tree/b85e06634a085b93155c455db2371c45ddf3ad2b/Week%2031

Notebook file direct link

<https://github.com/mramsundar/IITM-Graded->

Projects/blob/b85e06634a085b93155c455db2371c45ddf3ad2b/Week%2031/tweet_anaysis.ipynb

PDF of the notebook

<https://github.com/mramsundar/IITM-Graded->

Projects/blob/b85e06634a085b93155c455db2371c45ddf3ad2b/Week%2031/tweet_anaysis.pdf

Documentation

1. My approach started with defining a set of global constants, that could be then reused in multiple places across the notebook.

```
In [3]: @dataclass
class Config:
    DATA_PATH: str = "twitter_training.csv"
    RANDOM_SEED: int = 28
    SPACY_MODEL: str = "en_core_web_sm"
    TEXT_VECTOR_MAX_TOKENS:int = 5000
    TEXT_VECTOR_OUTPUT_MODE: str = "int"
    TEXT_VECTOR_OUTPUT_SEQUENCE_LENGTH:int =300
    TEXT_VECTOR_EMBEDDING_DIMENSION:int = 64 # Recommended value for tweets and short texts
    VECTORIZER: Any = None
    TOTAL_EPOCHS: int = 10
    BATCH_SIZE: int = 256
    EARLY_STOPPING_PATIENCE: int = 5
    LSTM_UNITS: int = 64 # Keeping the embedding dimension size in mind
    HIDDEN_UNITS: int = 32
    DROPOUT_RATE: float = 0.2
```

```
GlobalConfig = Config()

def set_global_seed(seed_value: int) -> None:
    os.environ['PYTHONHASHSEED'] = str(seed_value)
    random.seed(seed_value)
    tf.random.set_seed(seed_value)
    np.random.seed(seed_value)

set_global_seed(GlobalConfig.RANDOM_SEED)
```

2. Dataset was then loaded and basic cleanup process started, including setting headers, removing duplicates

```
def load_and_setup_data() -> pd.DataFrame:
    df = pd.read_csv(GlobalConfig.DATA_PATH)
    # First column appears to be some sort file or sequence number and the second appears to be the source
    # Those two columns do not have any impact on sentiment analysis
    # Take the last two columns
    df = df.iloc[:, -2:]
    # Swap columns 1 and 2
    df = df[[df.columns[1], df.columns[0]]]
    # Setup column names
    df.columns = ["tweet", "sentiment"]
    # Remove empty rows
    df = df.dropna()
    # Remove where sentiment is "Irrelevant"
    df = df[df["sentiment"] != "Irrelevant"]
    # Remove duplicate rows
    df = df.drop_duplicates()
    # Look for tweets where the same tweet is classified as a different sentiment
    # Take the first occurrence - this will get us clean data and will not mislead the classifier later during train
    df = df.drop_duplicates(subset=["tweet"], keep="first")
    return df
```

```
df = load_and_setup_data()
```

3. Since this is a sentiment analysis project, the next step was to take what is essential from the tweets. SpaCy was used to lemmatize the text

```
df = load_and_setup_data()

# Just keep the essentials, remove named entities, parsing and sentence segmentation for speed
nlp = spacy.load(GlobalConfig.SPACY_MODEL, disable=["ner", "parser", "senter"])

def clean_and_pre_process(texts: list[str]) -> list[str]:
    # Standard regex patterns for cleaning tweets
    url_pattern = re.compile(r"http\S+|www\S+|https\S+")
    twitter_handle_pattern = re.compile(r"@\w+\|\#\w+")
    non_alpha_pattern = re.compile(r"^\w+|[^\w\-\_]+")

    cleaned_texts: list[str] = []
    for text in texts:
        text = url_pattern.sub("", text)
        text = twitter_handle_pattern.sub("", text)
        text = non_alpha_pattern.sub("", text)
        cleaned_texts.append(text.strip().lower())

    # Use spaCy's nlp.pipe for efficient processing of multiple texts
    docs = nlp.pipe(cleaned_texts, batch_size=2000, n_process=-1)

    final_texts: list[str] = []
    for doc in tqdm(docs, total=len(cleaned_texts)):
        # This is a key line - we are lemmatizing, removing stop words, punctuation and single character tokens
        tokens = [token.lemma_ for token in doc if not token.is_stop and not token.is_punct and len(token.lemma_) > 1]
        final_texts.append(" ".join(tokens))

    return final_texts

df["sanitized_tweet"] = clean_and_pre_process(df["tweet"].tolist())
df.head()
```

4. Next is feature engineering. We transformed the cleaned-up tweet into a sequence of integers. These would later be converted into embeddings when we construct our RNN


```
def create_tokenized_words(df: pd.DataFrame) -> tuple[pd.DataFrame, Any]:
    texts = df["sanitized_tweet"].values
    vectorizer = tf.keras.layers.TextVectorization(
        max_tokens=GlobalConfig.TEXT_VECTOR_MAX_TOKENS,
        output_mode=GlobalConfig.TEXT_VECTOR_OUTPUT_MODE,
        output_sequence_length=GlobalConfig.TEXT_VECTOR_OUTPUT_SEQUENCE_LENGTH
    )
    vectorizer.adapt(texts)
    # Convert text to integer sequences - which will then be fed into the embedding layer
    tfidf_vectors = vectorizer(texts)
    # add the tokenized words as a new column
    df["sanitized_tweet_vector"] = list(tfidf_vectors.numpy())
    return (df, vectorizer)

df, GlobalConfig.VECTORIZER = create_tokenized_words(df)
df.head()
```




5. During EDA, it was observed that the sentiments were fairly well distributed

2.1 Basic Statistics

```
# Explore the distribution of tweet sentiments (e.g., how many positive, negative, and neutral tweets are there?)  
# Plot a simple bar chart  
sentiment_counts = df["sentiment"].value_counts()  
plt.figure(figsize=(10, 6))  
plt.bar(sentiment_counts.index, sentiment_counts.values.tolist())  
# Add total counts on top of each bar  
for i, count in enumerate(sentiment_counts.values):  
    plt.text(i, count + 100, str(count), ha="center")  
plt.xlabel("Sentiment")  
plt.ylabel("Number of Tweets")  
  
plt.title("Distribution of Tweet Sentiments")  
plt.show()
```



6. Bar charts with word frequencies and word clouds were created next




7. Word clouds

```
In [8]: def create_wordcloud() -> None:
    # Wordcloud for positive and negative tweets
    for sentiment in ["Positive", "Negative"]:
        subset = df[df["sentiment"] == sentiment]
        all_words = " ".join(subset["sanitized_tweet"])
        wordcloud = WordCloud(width=1000, height=400, background_color="white").generate(all_words)
```

```
plt.figure(figsize=(10, 10))
plt.imshow(wordcloud, interpolation="bilinear")
plt.axis("off")
plt.title(f"Word Cloud for {sentiment} Tweets")
plt.show()

create_wordcloud()
```




Word Cloud for Negative Tweets




- Last step of EDA was to look at the relationship between tweet length and the sentiment. It was observed that there is no obvious relationship

```
In [9]: def show_relationship_between_tweet_length_and_sentiment() -> None:
    # Relationship between tweet length and sentiment
    df["tweet_length"] = df["tweet"].apply(lambda x: len(x.split()))
    plt.figure(figsize=(10, 6))
    for sentiment in sentiment_categories:
        subset = df[df["sentiment"] == sentiment]
        plt.hist(subset["tweet_length"], bins=50, alpha=0.3, label=sentiment)
    plt.xlabel("Tweet Length (in words)")
    plt.ylabel("Number of Tweets")
    plt.title("Distribution of Tweet Lengths by Sentiment")
    plt.legend()
    plt.show()

show_relationship_between_tweet_length_and_sentiment()
```



9. The following is the RNN model architecture (created in Mermaid format based on the Python code) with ~400K trainable parameters



10. Model code

```
[]: def build_rnn_model() -> keras.Model:  
    model = keras.Sequential()  
  
    #Input layer - same as output sequence length of text vectorization  
    model.add(keras.Input(shape=(GlobalConfig.TEXT_VECTOR_OUTPUT_SEQUENCE_LENGTH,)))  
  
    # Embedding layer  
    model.add(  
        keras.layers.Embedding(  
            input_dim=GlobalConfig.TEXT_VECTOR_MAX_TOKENS,  
            output_dim=GlobalConfig.TEXT_VECTOR_EMBEDDING_DIMENSION  
        )  
    )  
  
    # LSTM Layer for analyzing sequences  
    model.add(keras.layers.Bidirectional(  
        keras.layers.LSTM(GlobalConfig.LSTM_UNITS, dropout=GlobalConfig.DROPOUT_RATE)  
    ))  
  
    # Fully connected hidden layer  
    model.add(keras.layers.Dense(GlobalConfig.HIDDEN_UNITS, activation="relu"))  
  
    # Output layer set to 3 classes - positive, negative, neutral  
    model.add(keras.layers.Dense(3, activation="softmax"))  
  
    model.compile(loss="sparse_categorical_crossentropy", optimizer="adam", metrics=["accuracy"])  
    return model
```

```
model = build_rnn_model()  
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
embedding (Embedding)	(None, 300, 64)	320000
bidirectional (Bidirectional)	(None, 128)	66048
dense (Dense)	(None, 32)	4128
dense_1 (Dense)	(None, 3)	99
<hr/>		
Total params: 390275 (1.49 MB)		
Trainable params: 390275 (1.49 MB)		
Non-trainable params: 0 (0.00 Byte)		

11. Before model evaluation, final data preparation was implemented to map sentiments into integers and getting X and y values

```

: # Add an extra column "sentiment_label" with numerical labels - encoding
def prepare_data_for_model(df: pd.DataFrame) -> pd.DataFrame:
    label_mapping = {"Negative": 0, "Neutral": 1, "Positive": 2}
    df["sentiment_label"] = df["sentiment"].map(label_mapping)
    return df

df = prepare_data_for_model(df)

def get_features_and_targets(df: pd.DataFrame) -> tuple[np.ndarray, np.ndarray]:
    X = np.array(list(df["sanitized_tweet_vector"].values))
    y = df["sentiment_label"].to_numpy(dtype=np.int32)
    return X, y

X, y = get_features_and_targets(df)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=GlobalConfig.RANDOM_SEED, str

# Stop training if validation loss does not improve for certain number of epochs
early_stopping = keras.callbacks.EarlyStopping(monitor="val_loss", patience=GlobalConfig.EARLY_STOPPING_PATIENCE, r

history = model.fit(
    X_train,
    y_train,
    validation_split=0.2,
    epochs=GlobalConfig.TOTAL_EPOCHS,
    batch_size=GlobalConfig.BATCH_SIZE,
    callbacks=[early_stopping, TqdmCallback(verbose=1)], #TqdmCallback for progress bar
    verbose=0
)

100%|██████████| 10/10 [01:35<00:00,  9.55s/epoch, loss=0.285, accuracy=0.895, val_loss=0.562, val_accuracy=0.817]

```

12. Key points on model training

- a. Early stopping was implemented to ensure model does not overfit
- b. 10 Epochs were chosen based on industry standard practices for twitter sentiment analysis recommendations
- c. ~82% val_accuracy was obtained after this round of training

13. Model evaluation was done by analyzing training and validation loss on the one hand and training and validation accuracy on the other. Through this, Epoch 7 was identified as the right number as it gave the lowest validation loss

3.3 Evaluation

```
: def plot_training_history(history):
    metrics = history.history

    # Get key metrics
    loss = metrics["loss"]
    val_loss = metrics["val_loss"]
    accuracy = metrics["accuracy"]
    val_accuracy = metrics["val_accuracy"]
    epochs = range(1, len(loss) + 1)

    min_val_loss = min(val_loss)
    min_val_loss_epoch = val_loss.index(min_val_loss) + 1

    plt.figure(figsize=(14, 5))
    plt.subplot(1, 2, 1)
    plt.plot(epochs, loss, "bo-", label="Training Loss") #bo- Blue line with o markers and - line
    plt.plot(epochs, val_loss, "ro-", label="Validation Loss")


    plt.plot(min_val_loss_epoch, min_val_loss, 'gs', markersize=10, label='Lowest Val Loss') # Green Square
    plt.annotate(
        f"Lowest: {min_val_loss:.4f}",
        xy=(min_val_loss_epoch, min_val_loss),
        xytext=(min_val_loss_epoch, min_val_loss + 0.05), # Offset text slightly up
        horizontalalignment='center'
    )

    plt.title("Training and Validation Loss")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend()
    plt.grid(True)

    plt.subplot(1, 2, 2)
    plt.plot(epochs, accuracy, "bo-", label="Training Accuracy")
    plt.plot(epochs, val_accuracy, "ro-", label="Validation Accuracy")
    plt.title("Training and Validation Accuracy")
    plt.xlabel("Epochs")
    plt.ylabel("Accuracy")
    plt.legend()
    plt.grid(True)

    plt.tight_layout()
    plt.show()
```

```
plot_training_history(history)
```



14. To make things more interesting, we went ahead and completed all 10 Epochs and performed a what-if analysis

```
model_unrestricted = build_rnn_model()

# We are trying to do two things below - do not stop training early and train for full number of epochs and then do
print(f"Starting unrestricted training for {GlobalConfig.TOTAL_EPOCHS} epochs ...")
history_unrestricted = model_unrestricted.fit(
    X_train,
    y_train,
    validation_split=0.2,
    epochs=GlobalConfig.TOTAL_EPOCHS,
    batch_size=GlobalConfig.BATCH_SIZE,
    callbacks=[TqdmCallback(verbose=1)],
    verbose=0
)

Starting unrestricted training for 10 epochs ...
100%|██████████| 10/10 [01:35<00:00,  9.60s/epoch, loss=0.277, accuracy=0.897, val_loss=0.537, val_accuracy=0.824]
```

```
def plot_what_if_analysis_for_unrestricted_training(history):
    val_loss = history_unrestricted.history["val_loss"]
    epochs = range(1, len(val_loss) + 1)

    # Find the epoch where val_loss was minimum (Best Model)
    best_epoch_index = val_loss.index(min(val_loss))
    best_epoch = best_epoch_index + 1
    best_loss = val_loss[best_epoch_index]

    plt.figure(figsize=(10, 8))

    # Plot the curve
    plt.plot(epochs, val_loss, "r-o", label="Validation Loss", linewidth=2)

    # Vertical line at the Best Epoch
    plt.axvline(x=best_epoch, color="green", linestyle="--", label=f"Ideal Stop (Epoch {best_epoch})")


    # Shade the Overfitting Zone (What happened after)
    if best_epoch < len(epochs):
        plt.axvspan(best_epoch, len(epochs), color="red", alpha=0.1, label="Overfitting / Wasted Compute")

    # Add annotation arrow
    plt.annotate(
        "Loss increases here\n(Model gets worse)",
        xy=(len(epochs), val_loss[-1]),
        xytext=(best_epoch + 2, max(val_loss) * 0.95)
    )

    plt.title("What-If Analysis: Training Without Early Stopping")
    plt.xlabel("Epochs")
    plt.ylabel("Validation Loss")
    plt.legend()
    plt.grid(True)
    plt.show()

# Run the plot
plot_what_if_analysis_for_unrestricted_training(history_unrestricted)
```

15. This also confirmed our earlier finding that beyond Epoch 7, our returns were diminishing



16. Model Cross Validation was then performed to find out how model behaves for un-seen data outside the validation dataset. This also yielded ~83% validation accuracy – further attesting the fact that our training is very solid

3.4 Model Improvement

```
In [ ]: # Evaluate the model to see if our model is performing well on unseen test data through fold cross validation
def do_model_cross_validation() -> None:
    skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=GlobalConfig.RANDOM_SEED)
    fold_no = 1
    accuracies = []

    for train_index, val_index in skf.split(X, y):
        X_train_fold, X_val_fold = X[train_index], X[val_index]
        y_train_fold, y_val_fold = y[train_index], y[val_index]

        model_cv = build_rnn_model()

        history_cv = model_cv.fit(
            X_train_fold,
            y_train_fold,
            validation_data=(X_val_fold, y_val_fold),
            epochs=GlobalConfig.TOTAL_EPOCHS,
            batch_size=GlobalConfig.BATCH_SIZE,
            callbacks=[early_stopping],
            verbose=0
        )

        scores = model_cv.evaluate(X_val_fold, y_val_fold, verbose=0)
        print(f"Fold {fold_no} - Validation Accuracy: {scores[1]*100:.2f}%")
        accuracies.append(scores[1] * 100)
        fold_no += 1

    print(f"Average Validation Accuracy across folds: {np.mean(accuracies):.2f}% ± {np.std(accuracies):.2f}%")

do_model_cross_validation()
```

```
Fold 1 - Validation Accuracy: 83.90%
Fold 2 - Validation Accuracy: 83.70%
Fold 3 - Validation Accuracy: 83.20%
Fold 4 - Validation Accuracy: 82.98%
Fold 5 - Validation Accuracy: 83.34%
Average Validation Accuracy across folds: 83.42% ± 0.33%
```

-
17. For tuning the hyper parameters, Grid Search was performed next to determine the right combination of LSTM, Hidden layer units and Dropout rates. 64 was kept as the embedding dimension, the LSTM and hidden layer units to align with industry standard recommendations in this space. 0.2 was identified as the optimal dropout rate yielding ~81% accuracy

```

: def do_model_grid_search(lstm_units_list: list[int], hidden_units_list: list[int], dropout_list: list[float]) -> No
    best_accuracy = 0.0
    best_params = {}
    # Loop through all possible combinations
    for lstm_unit in lstm_units_list:
        for hidden_unit in hidden_units_list:
            for dropout in dropout_list:
                print(f"Training model with {lstm_unit} lstm units, {hidden_unit} hidden units, and {dropout} dropout")
                model_gs = keras.Sequential()
                model_gs.add(keras.layers.Input(shape=(GlobalConfig.TEXT_VECTOR_OUTPUT_SEQUENCE_LENGTH,)))
                model_gs.add(
                    keras.layers.Embedding(
                        input_dim=GlobalConfig.TEXT_VECTOR_MAX_TOKENS,
                        output_dim=GlobalConfig.TEXT_VECTOR_EMBEDDING_DIMENSION
                    )
                )
                model_gs.add(keras.layers.Bidirectional(
                    keras.layers.LSTM(lstm_unit, dropout=dropout)
                ))
                model_gs.add(keras.layers.Dense(hidden_unit, activation="relu"))

                model_gs.add(keras.layers.Dense(3, activation="softmax"))
                model_gs.compile(loss="sparse_categorical_crossentropy", optimizer="adam", metrics=["accuracy"])

                history_gs = model_gs.fit(
                    X_train,
                    y_train,
                    validation_split=0.2,
                    epochs=10,
                    batch_size=GlobalConfig.BATCH_SIZE,
                    callbacks=[early_stopping],
                    verbose=0
                )
                # Capture scores and get best accuracy

```

```

scores = model_gs.evaluate(X_test, y_test, verbose=0)
accuracy = scores[1] * 100
print(f"Test Accuracy: {accuracy:.2f}%")

if accuracy > best_accuracy:
    best_accuracy = accuracy
    best_params = {"lstm unit": lstm_unit, "hidden unit": hidden_unit, "dropout": dropout}

print(f"Best Test Accuracy: {best_accuracy:.2f}% with parameters: {best_params}")

do_model_grid_search(lstm_units_list=[64], hidden_units_list=[64], dropout_list=[0.0, 0.2, 0.3])

Training model with 64 lstm units, 64 hidden units, and 0.0 dropout...
Test Accuracy: 79.48%
Training model with 64 lstm units, 64 hidden units, and 0.2 dropout...
Test Accuracy: 81.02%
Training model with 64 lstm units, 64 hidden units, and 0.3 dropout...
Test Accuracy: 80.99%
Best Test Accuracy: 81.02% with parameters: {'lstm unit': 64, 'hidden unit': 64, 'dropout': 0.2}

```

18. A benchmark test was performed with twitter data from NTLK to test our model strength. The outcome of the test was moderate at best. This indicates that our training data is fairly domain specific and needs to be utilizing transfer learning by taking into consideration a wider vocabulary

Benchmarking our Model with external data and making the case for transfer learning

```
5]: def quick_cleanse(text: str) -> str:
    if not isinstance(text, str):
        return ""
    text = html.unescape(text)
    text = text.lower()
    text = re.sub(r"http\S+|www\S+|https\S+", "", text, flags=re.MULTILINE)
    text = re.sub(r"@w+", "", text)
    text = re.sub(r"#", "", text)
    text = re.sub(r"[a-zA-Z\s]", "", text)
    text = re.sub(r"\s+", " ", text).strip()
    return text

def benchmark_model():
```

```
print("Downloading NLTK Twitter dataset...")
nltk.download("twitter_samples", quiet=True)
raw_pos = [quick_cleanse(t) for t in twitter_samples.strings("positive_tweets.json")[:1000]]
raw_neg = [quick_cleanse(t) for t in twitter_samples.strings("negative_tweets.json")[:1000]]

X_benchmark_text = raw_pos + raw_neg
y_benchmark_true = np.array([2] * 1000 + [0] * 1000)

print(f"Vectorizing {len(X_benchmark_text)} tweets...")
X_benchmark_vectors = GlobalConfig.VECTORIZER(X_benchmark_text) # Reusing the saved vectorizer

# 5. Predict with Your Model
print("Running predictions...")
y_pred_probs = model.predict(X_benchmark_vectors, verbose=0)
y_pred_probs[:, 1] = -1 # No neutral class in benchmark dataset
y_pred_classes = np.argmax(y_pred_probs, axis=1)

print("\n" + "="*40)
print(" EXTERNAL DATASET BENCHMARK RESULTS")
print("="*40)

print(classification_report(
    y_benchmark_true,
    y_pred_classes,
    labels=[0, 1, 2],
    target_names=["Negative", "Neutral", "Positive"]
))

cm = confusion_matrix(y_benchmark_true, y_pred_classes, labels=[0, 1, 2])

fig, ax = plt.subplots(figsize=(8, 6))
cax = ax.matshow(cm, cmap="Blues")
fig.colorbar(cax)

ax.set_xticklabels([""] + ["Pred Negative", "Pred Neutral", "Pred Positive"])
ax.set_yticklabels([""] + ["True Negative", "True Neutral", "True Positive"])
```

```
ax.xaxis.set_ticks_position("bottom")


for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, str(cm[i, j]), va="center", ha="center", color="black", fontsize=12)

plt.title("Confusion Matrix: My Model vs. NLTK Dataset")
plt.xlabel("Predicted Label")
plt.ylabel("Actual Label")
plt.show()

benchmark_model()
```

Downloading NLTK Twitter dataset

EXTERNAL DATASET BENCHMARK RESULTS				
	precision	recall	f1-score	support
Negative	0.54	0.33	0.41	1000
Neutral	0.00	0.00	0.00	0
Positive	0.52	0.72	0.60	1000
accuracy			0.53	2000
macro avg	0.35	0.35	0.34	2000
weighted avg	0.53	0.53	0.51	2000



Observations

- Our accuracy with random tweets is as good as a coin toss
- This indicates that our vocabulary from `twitter_training.csv` is limited
- We could use a model like GloVe embeddings through transfer learning method

19. Finally, a quick test with sample tweets were performed and the model performed as expected

```
def manual_test_samples() -> pd.DataFrame:
    sample_tweets = [
        "I absolutely love the new features!",
        "The service was terrible and I am very disappointed.",
        "It was okay, nothing special but not bad either.",
        "Can anyone help me with this issue? I'm stuck."
    ]
    label_mapping = {"Negative": 0, "Neutral": 1, "Positive": 2}
    index_to_label = {v: k for k, v in label_mapping.items()}
    processed_samples = GlobalConfig.VECTORIZER(sample_tweets)
    predictions = model.predict(processed_samples)
    predicted_indices = np.argmax(predictions, axis=1)
    predicted_labels = [index_to_label[idx] for idx in predicted_indices]
    results_df = pd.DataFrame({
        "Tweet": sample_tweets,
        "Prediction": predicted_labels,
        "Confidence": np.max(predictions, axis=1)
    })

    return results_df

results_df = manual_test_samples()
display(results_df)
```

✓ 0.0s

1/1 [=====] - 0s 21ms/step

	Tweet	Prediction	Confidence
0	I absolutely love the new features!	Positive	0.837939
1	The service was terrible and I am very disappo...	Negative	0.963436
2	It was okay, nothing special but not bad either.	Neutral	0.723399
3	Can anyone help me with this issue? I'm stuck.	Negative	0.918900

Presentation

Presentation is located here in GitHub - <https://github.com/mramsundar/IITM-Graded-Projects/blob/c623766d58d22acec56bc2c02f906e3c3aaf0634/Week%2031/Week%2031%20Summary%20Presentation.pdf>