

# WEEK 31 ASSIGNMENT SUMMARY

RAM SUNDAR

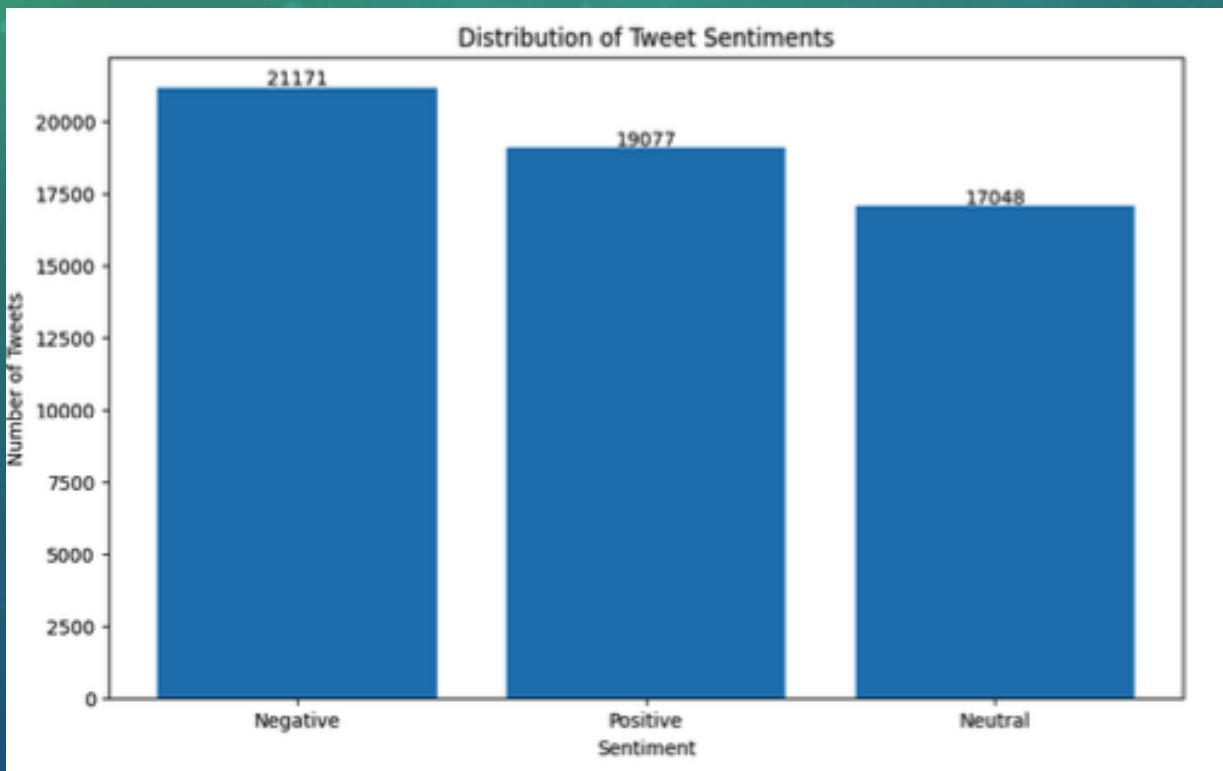
# OBJECTIVES

- Understand the structure of the Twitter dataset
- Perform data preprocessing and text cleaning
- Conduct exploratory data analysis to uncover insights
- Build a Recurrent Neural Network (RNN) model for sentiment classification
- Evaluate and improve the performance of the model
- Present findings and recommendations

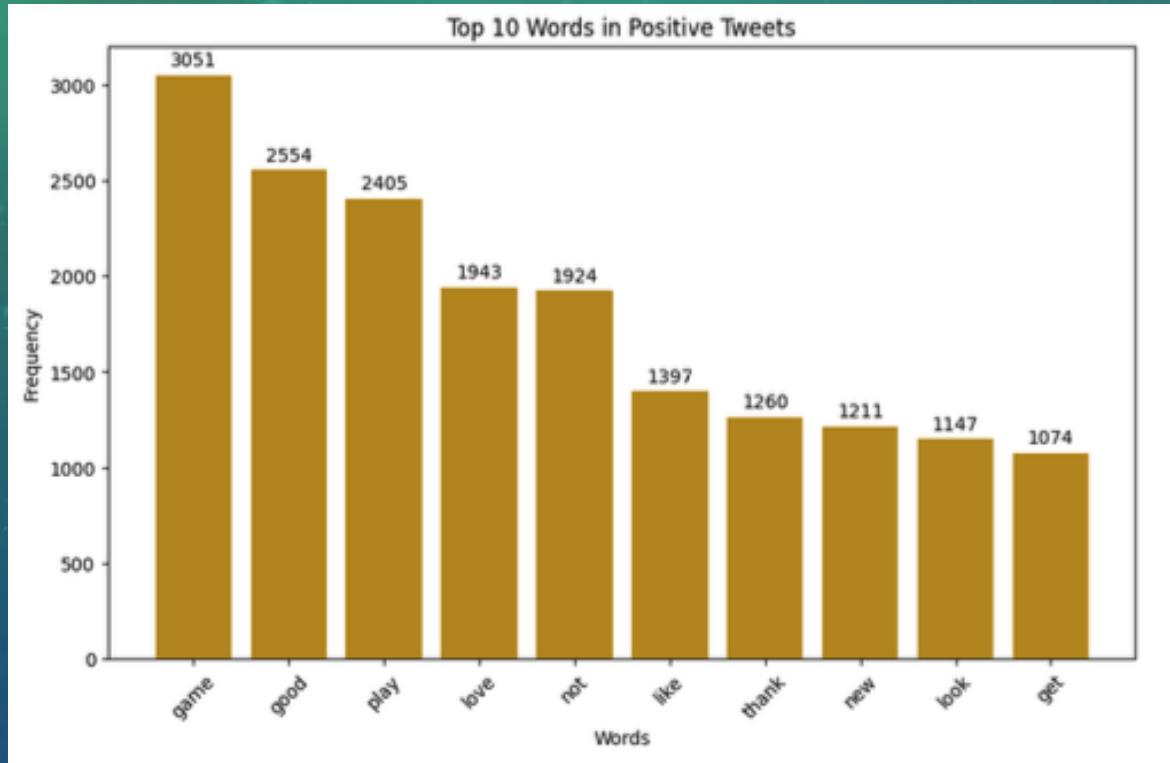
# DATASET OVERVIEW

- Standard CSV without headers
- Appears to be from multiple sources
- Dataset observations
  - Total Rows: 74681
  - Total Columns: 4
  - Duplicate Rows: 2700 (3.62%)
  - Total Missing Values: 686
  - Completely Empty Rows: 0
- Dataset after post processing
  - Total Rows: 57296
  - Total Columns: 2
  - Duplicate Rows: 0 (0.00%)
  - Total Missing Values: 0
  - Completely Empty Rows: 0

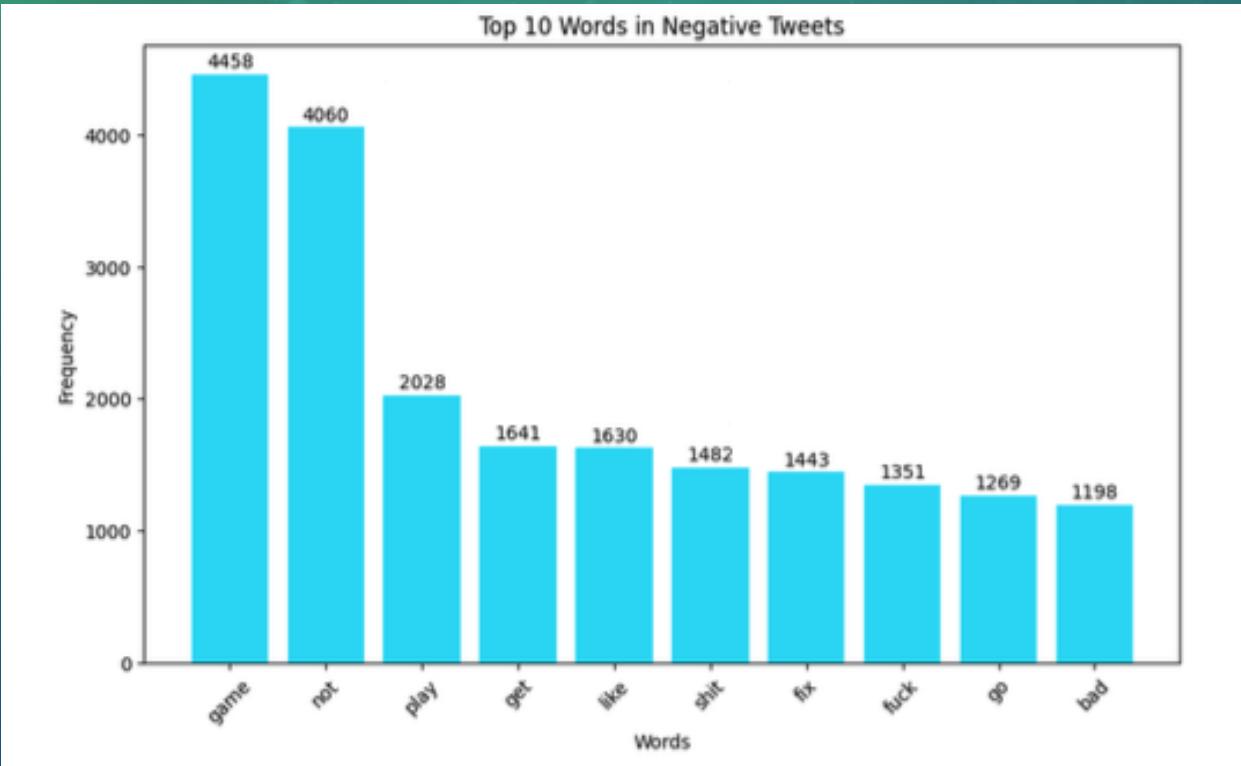
# EDA – KEY FINDINGS – SENTIMENT DISTRIBUTION



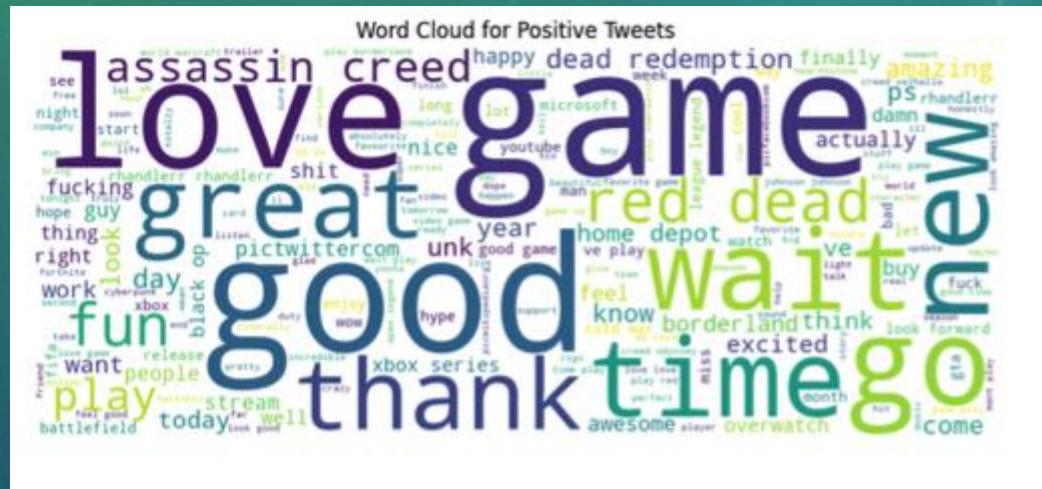
# EDA – KEY FINDINGS – TOP 10 POSITIVE KEYWORDS



# EDA – KEY FINDINGS – TOP 10 NEGATIVE KEYWORDS



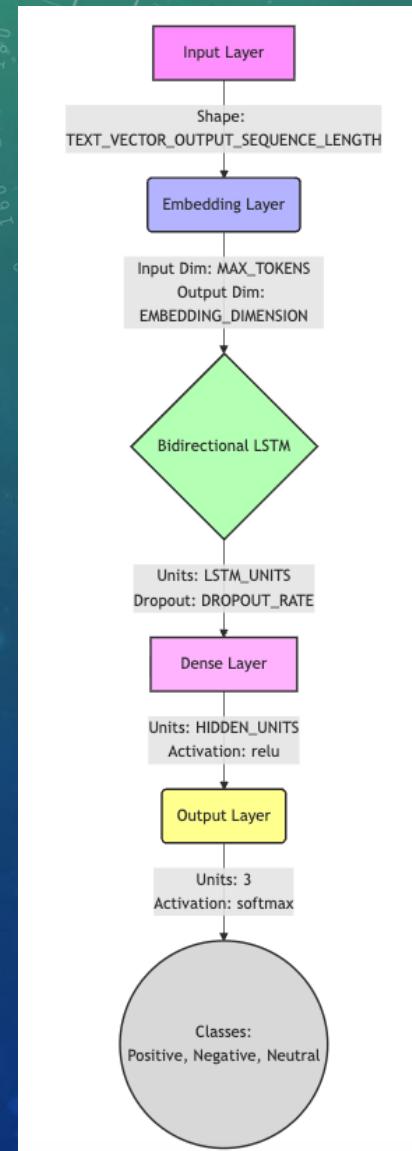
# EDA – KEY FINDINGS – WORD CLOUDS



# RNN MODEL BUILD - METHODOLOGY

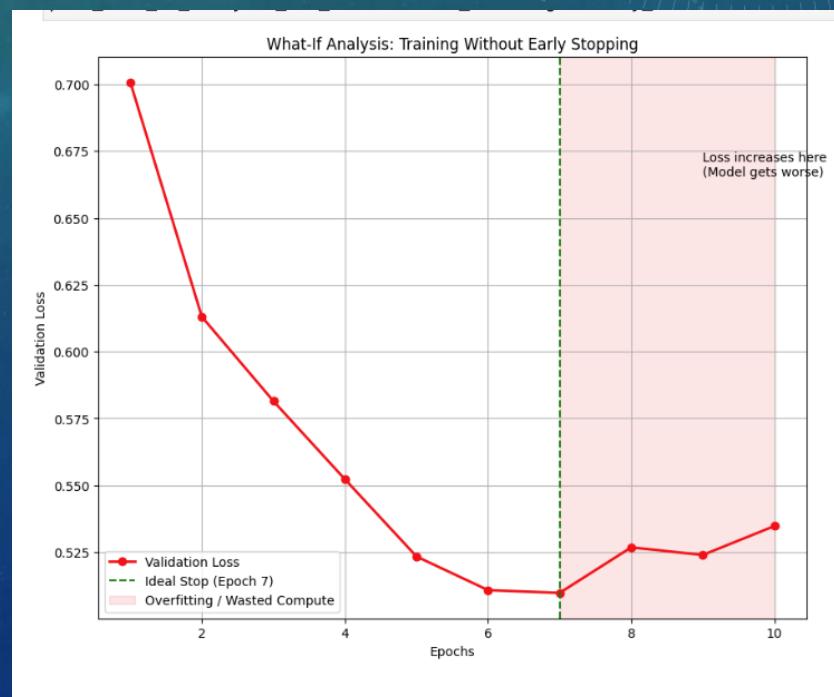
- Deep NN was chosen with multiple layers
- Embedding layer to begin with
- Bi Directional LSTM
- Followed by Dense Relu
- Lastly softmax with 3 classes for prediction probabilities

Model: "sequential"		
Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 300, 64)	320000
bidirectional (Bidirectional)	(None, 128)	66048
dense (Dense)	(None, 32)	4128
dense_1 (Dense)	(None, 3)	99
Total params: 390275 (1.49 MB)		
Trainable params: 390275 (1.49 MB)		
Non-trainable params: 0 (0.00 Byte)		



# EVAL RESULTS AND PERFORMANCE METRICS

- 10 Epochs were chosen – as per industry standard recommendation for tweet sentiment analysis
- Obtained ~82% Accuracy
- Performed what-if analysis to find out impact of not stopping early



# CHALLENGES FACED & MODEL IMPROVEMENT

- Cross validation was performed with 5 folds
- Model performance was consistent with an average validation accuracy of 83.5%
- Grid search with various dropout rates were performed as well
- Our RNN model and hyperparameter choice was proven solid
- External benchmarking was performed with NLTK twitter dataset
- Results were moderate at best, indicating the need to use a global vocabulary for embeddings generation

Fold 1 - Validation Accuracy: 82.08%

Fold 2 - Validation Accuracy: 83.73%

Fold 3 - Validation Accuracy: 83.04%

Fold 4 - Validation Accuracy: 83.02%

Fold 5 - Validation Accuracy: 83.37%

Average Validation Accuracy across folds: 83.05% ± 0.55%

Training model with 64 lstm units, 64 hidden units, and 0.0 dropout...

Test Accuracy: 79.41%

Training model with 64 lstm units, 64 hidden units, and 0.2 dropout...

Test Accuracy: 81.18%

Downloading NLTK Twitter dataset...

Vectorizing 2000 tweets...

Running predictions...

## ===== EXTERNAL DATASET BENCHMARK RESULTS =====

	precision	recall	f1-score	support
Negative	0.54	0.33	0.41	1000
Neutral	0.00	0.00	0.00	0
Positive	0.52	0.72	0.60	1000
accuracy			0.53	2000
macro avg	0.35	0.35	0.34	2000
weighted avg	0.53	0.53	0.51	2000

# SAMPLE TWEET DEMONSTRATION

- Model performed very well when tested with sample tweets – accurately predicting all inputs given

```
def manual_test_samples() -> pd.DataFrame:  
    sample_tweets = [  
        "I absolutely love the new features!",  
        "The service was terrible and I am very disappointed.",  
        "It was okay, nothing special but not bad either.",  
        "Can anyone help me with this issue? I'm stuck."  
    ]  
    label_mapping = {"Negative": 0, "Neutral": 1, "Positive": 2}  
    index_to_label = {v: k for k, v in label_mapping.items()}  
    processed_samples = GlobalConfig.VECTORIZER(sample_tweets)  
    predictions = model.predict(processed_samples)  
    predicted_indices = np.argmax(predictions, axis=1)  
    predicted_labels = [index_to_label[idx] for idx in predicted_indices]  
    results_df = pd.DataFrame({  
        "Tweet": sample_tweets,  
        "Prediction": predicted_labels,  
        "Confidence": np.max(predictions, axis=1)  
    })
```

```
return results_df  
  
results_df = manual_test_samples()  
display(results_df)
```

	Tweet	Prediction	Confidence
0	I absolutely love the new features!	Positive	0.837939
1	The service was terrible and I am very disappo...	Negative	0.963436
2	It was okay, nothing special but not bad either.	Neutral	0.723399
3	Can anyone help me with this issue? I'm stuck.	Negative	0.918900

